

# Goal Today: Different Approaches to Hyperparameter Tuning

How can we tune hyperparameters?

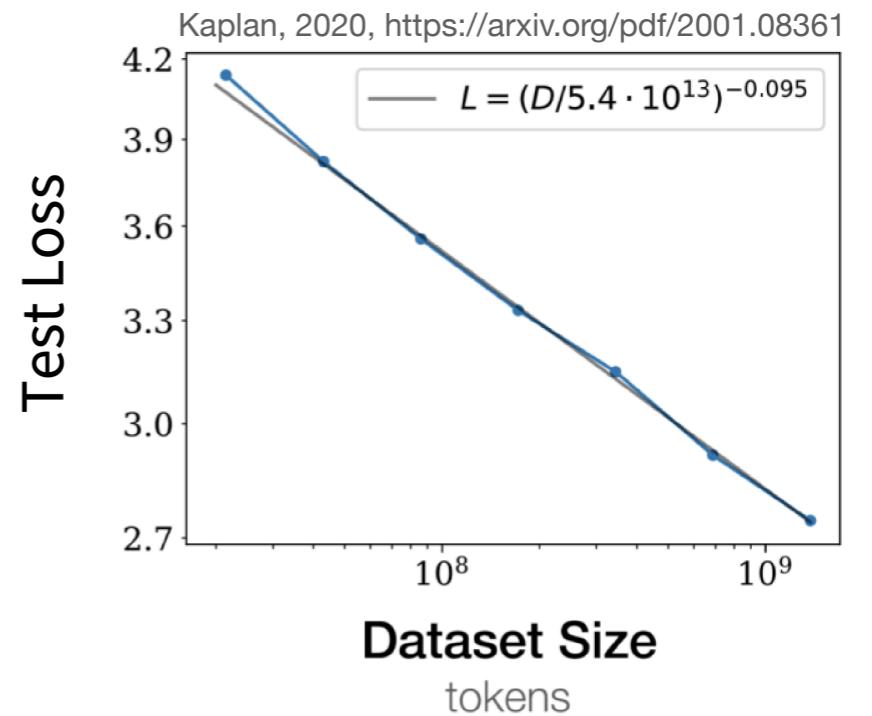
- Try and Pray
- Grid Search (costly)
- Do small-scale experiments. Then “extrapolate”

## 1. Draw a line: Scaling Law

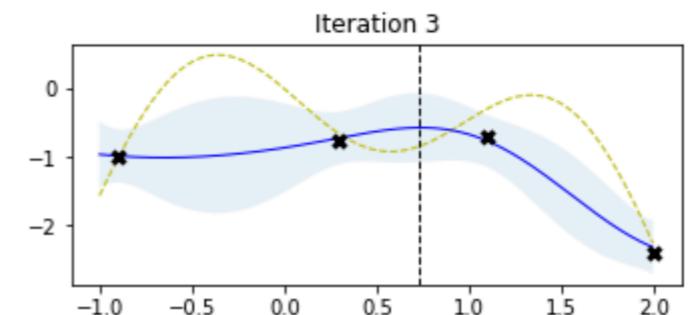
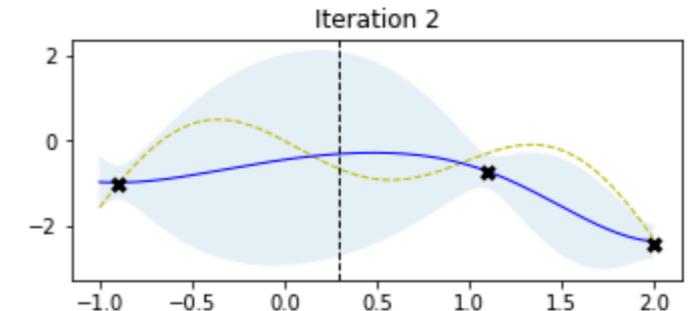
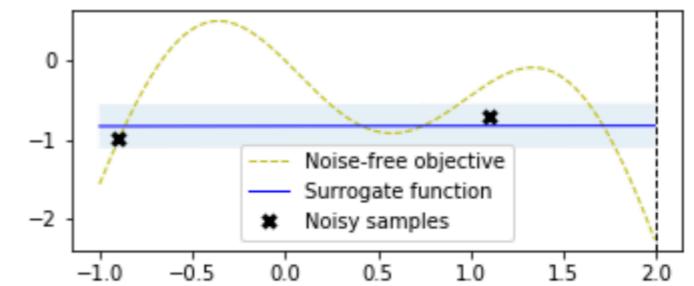
<https://stanford-cs324.github.io/winter2022/assets/pdfs/Scaling%20laws%20pdf.pdf>

## 2. (Multi-fidelity) Bayesian Optimization

## 3. Update hyperparameters online (specifically data)



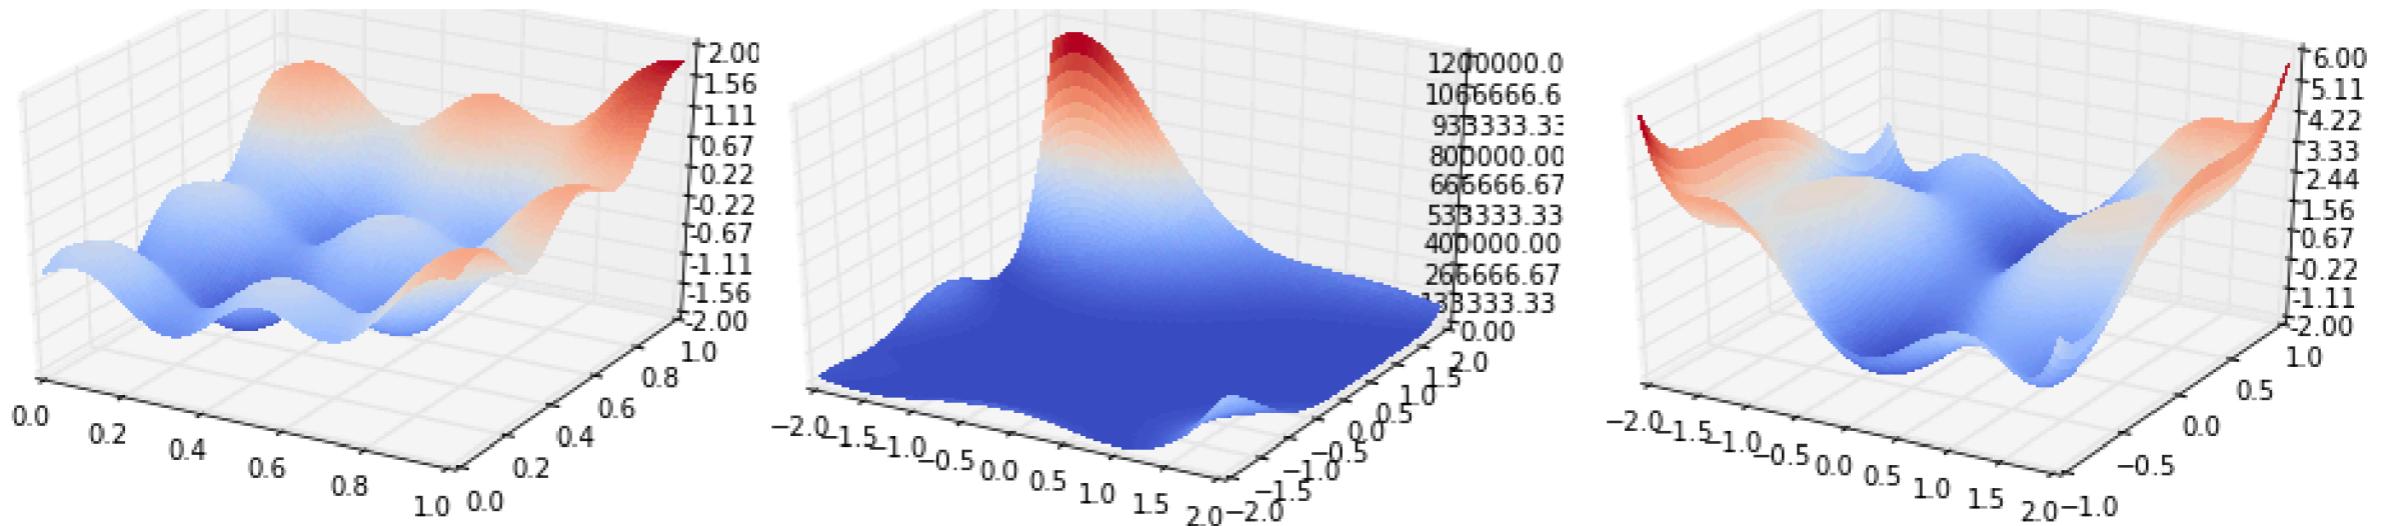
<https://krasserm.github.io/2018/03/21/bayesian-optimization/>



# Hyperparameter Optimization

Problem: which training parameters should I use?

- 0.01 or 0.001 learning rate?
- 0.9 or 0.99 momentum?
- 4 or 5 decoder blocks?



Evaluation of  $f$  is expensive

# Hyperparameter Tuning is Costly



Consumption	CO <sub>2</sub> e (lbs)
Air travel, 1 passenger, NY↔SF	1984
Human life, avg, 1 year	11,023
American life, avg, 1 year	36,156
Car, avg incl. fuel, 1 lifetime	126,000

Training one model (GPU)	
NLP pipeline (parsing, SRL)	39
w/ tuning & experimentation	78,468
Transformer (big)	192
w/ neural architecture search	626,155

Table 1: Estimated CO<sub>2</sub> emissions from training common NLP models, compared to familiar consumption.<sup>1</sup>

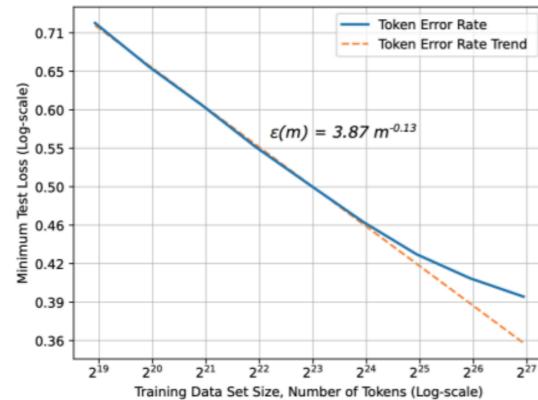


GPT-4 cost more than **100 Millions!**

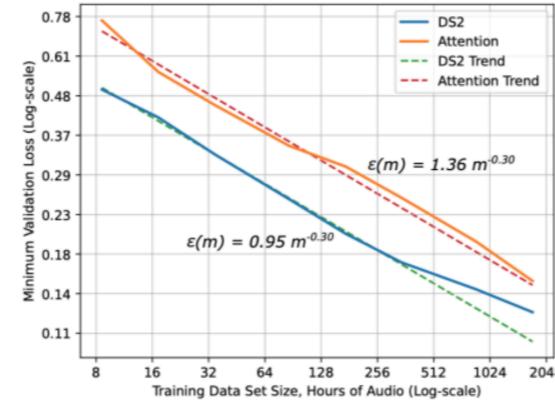
Infeasible to train many models

# Scaling Laws: An Interesting Phenomenon

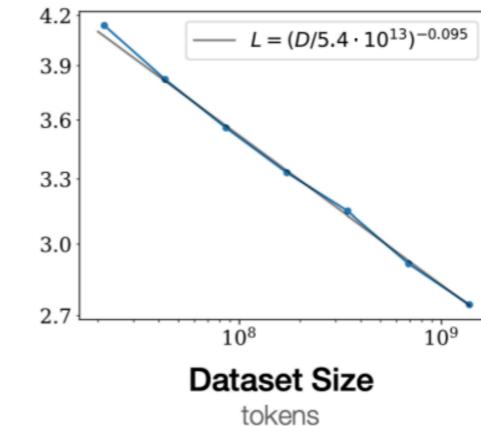
Scaling laws hold in many domains



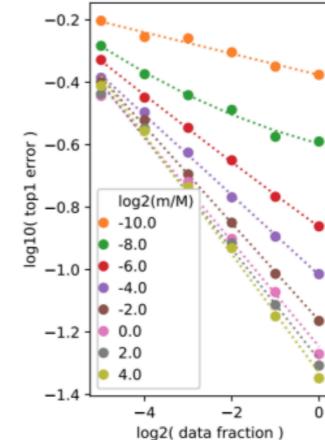
Machine translation



Speech



Language modeling



Object recognition

Hestness et al 2017.

Kaplan et al 2020.

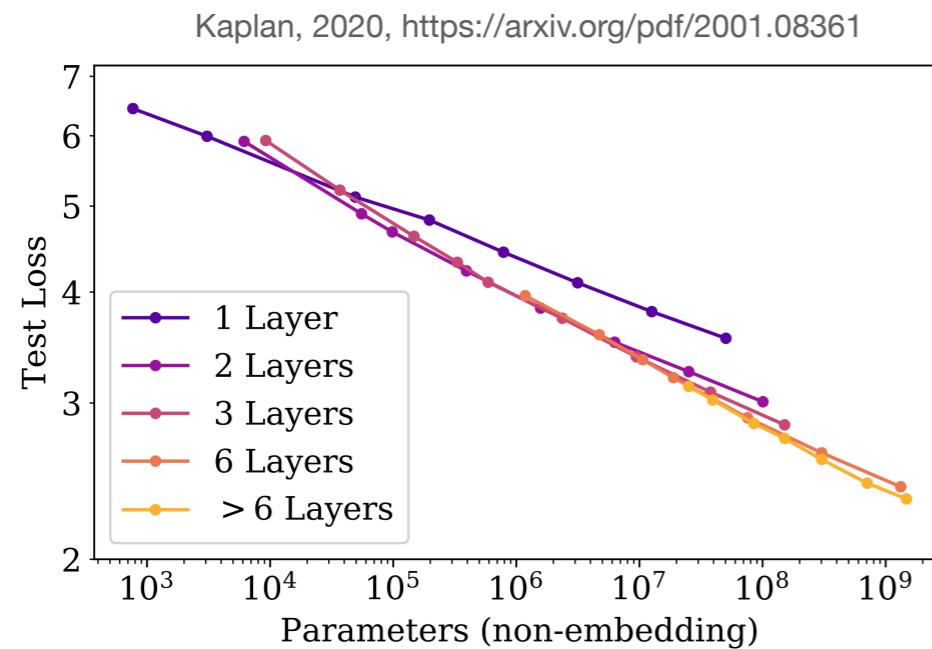
Rosenfeld 2020.

Significance:

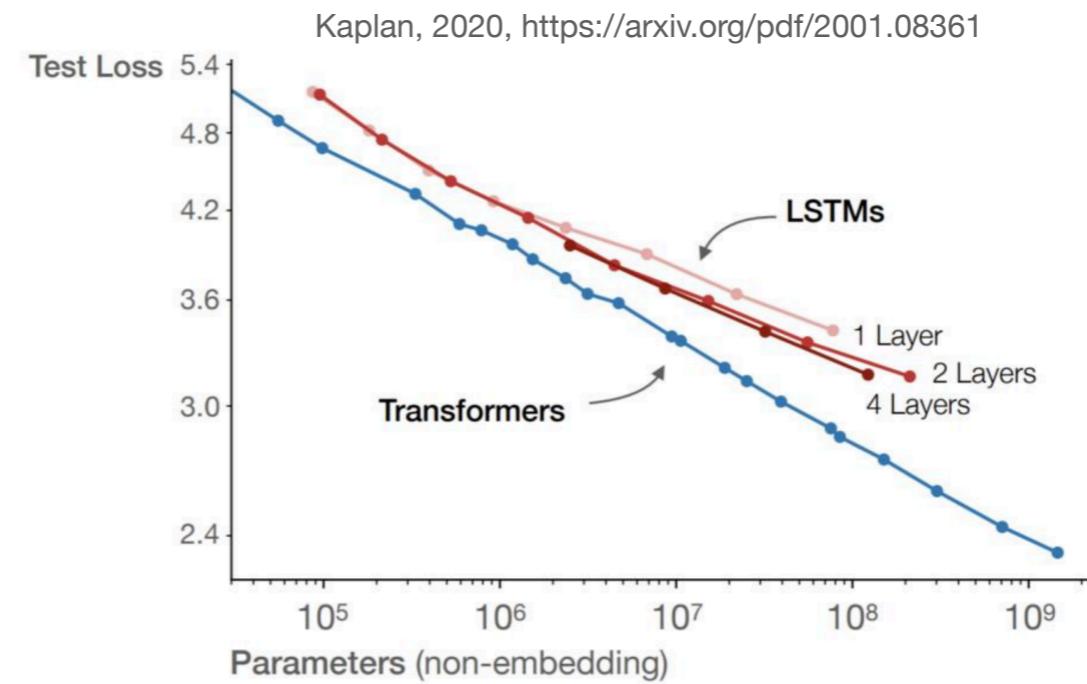
1. Predicting large-scale results => Efficient Design Choice
2. Allow low-budget contribution from research community
3. Enable resource allocation decisions (# of nuclear plants/gpus needed)

# Scaling Law: Identify Best Hyperparameters From Trend

Best # of layers?

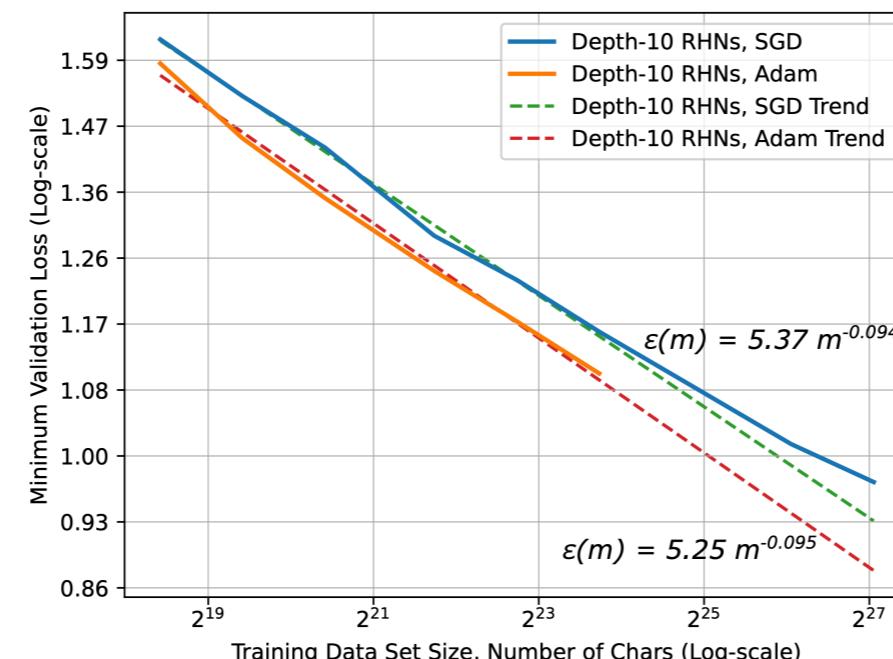


Best architecture?



Best optimizer?

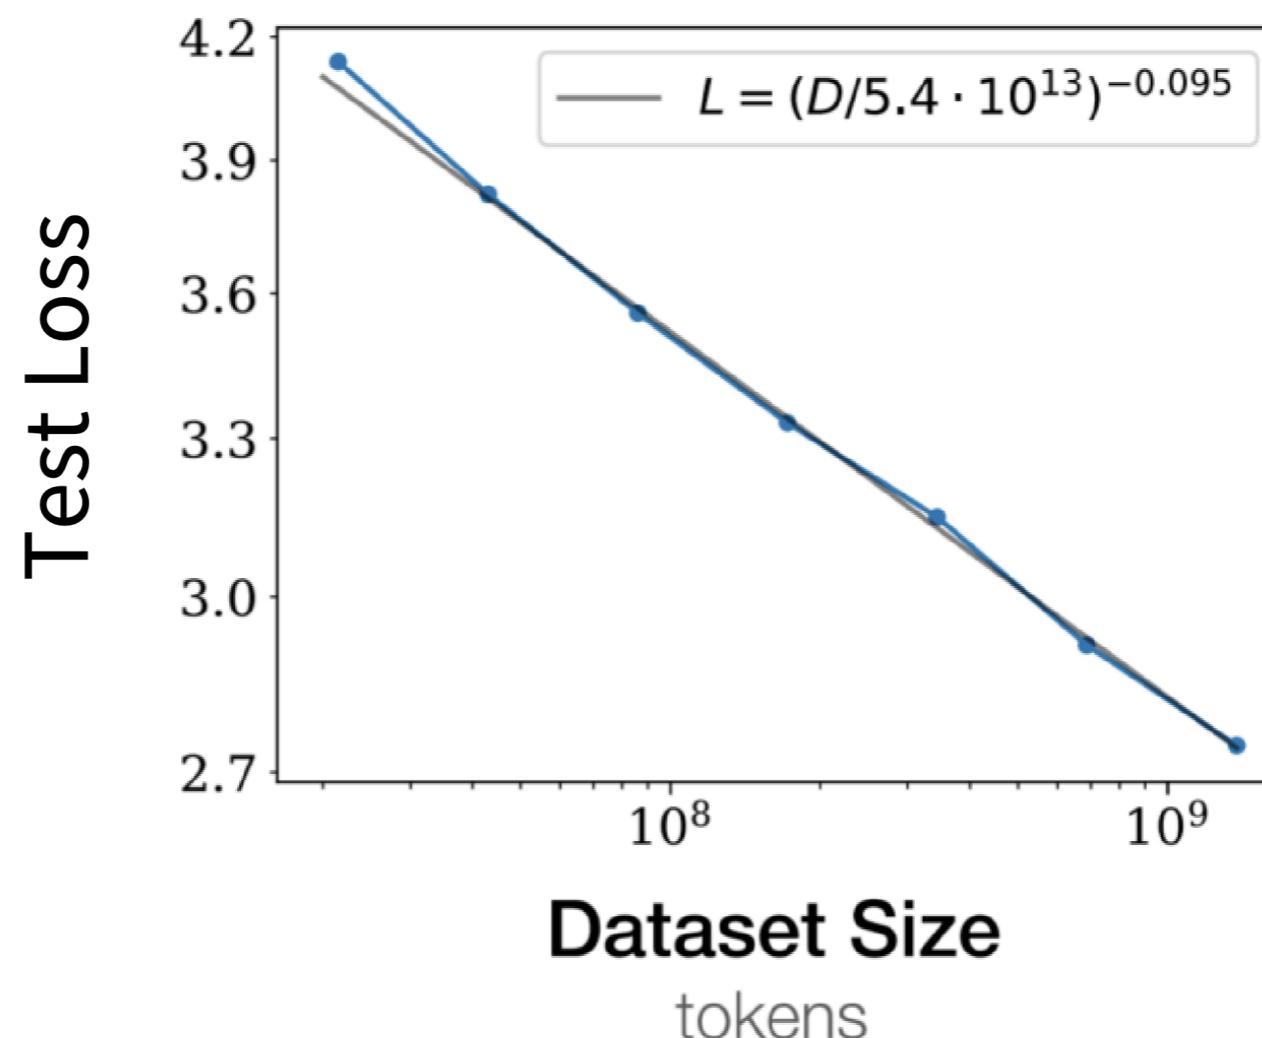
Hestness, 2017, <https://arxiv.org/pdf/1712.00409>



# Data Scaling Laws for Language Models

An empirical observation:

**Loss and dataset size is linear on a log-log plot**



Kaplan, 2020, <https://arxiv.org/pdf/2001.08361>

<https://stanford-cs324.github.io>

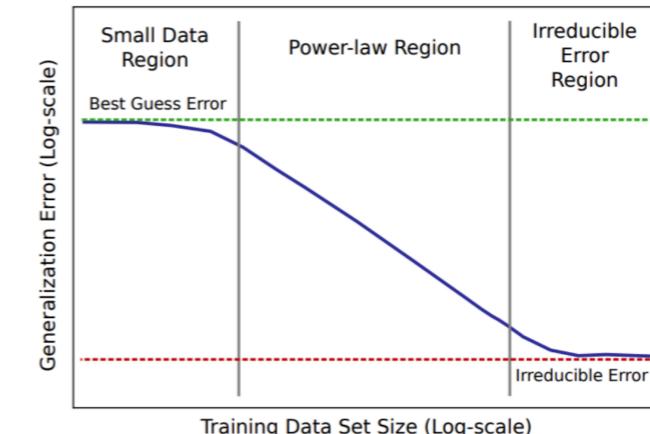
# Conceptual foundations of data scaling laws.

**Q:** Why do scaling laws show up?

We know error should be monotone



But why is it a power law / linear in log-log?



**A:** Estimation error naturally decays polynomially.

But this answer may take a moment to understand. Let's work through an example.

**Example:** If our task is to estimate the mean of a dataset, what's the scaling law?

# Toy example: mean estimation

**Input:**  $x_1 \dots x_n \sim N(\mu, \sigma^2)$

**Task:** estimate the average as  $\hat{\mu} = \frac{\sum_i x_i}{n}$

**What's the error?** By standard arguments..

$$E[(\hat{\mu} - \mu)^2] = \frac{\sigma^2}{n}$$

**This is a scaling law!!**

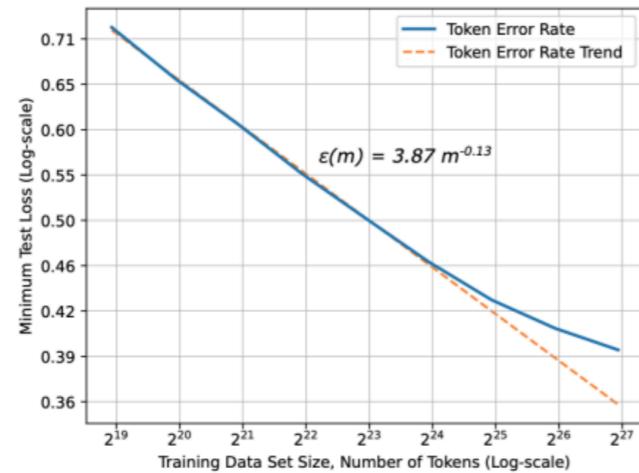
$$\log(\text{Error}) = -\log n + 2 \log \sigma$$

More generally, any polynomial rate  $1/n^\alpha$  is a scaling law

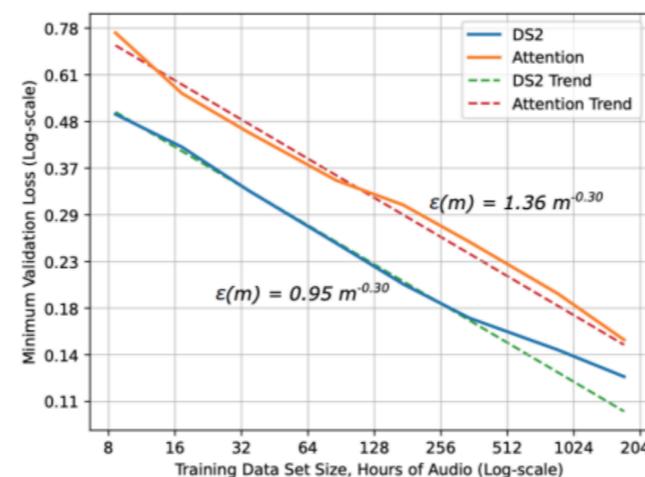
# Scaling law exponents: an intriguing mystery

**Fact:** Similar arguments show most ‘classical’ models (regression, etc) have  $\frac{1}{n}$  scaling

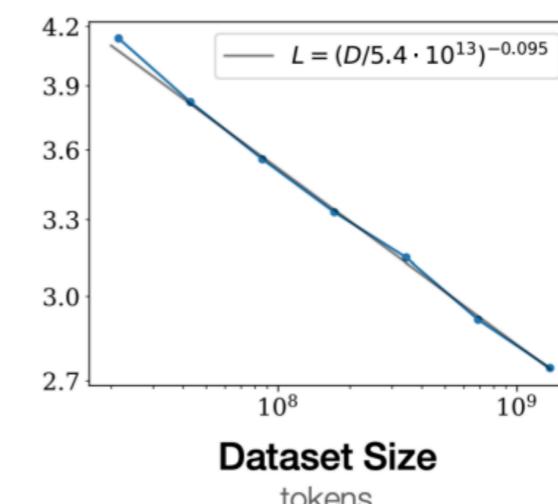
This means we should see  $y = -x + C$   
What do we find in neural scaling laws?



Machine translation



Speech



Language modeling

Very different from predictions.. Why might this be?

# Detour: scaling laws for (nonparametric) learning

Neural nets can approximate arbitrary functions. Lets turn that into an example.

**Input:**  $x_1 \dots x_n$  uniform in 2D unit box.  $y_i = f(x_i) + N(0,1)$

**Task:** estimate  $f(x)$

**Approach:** cut up the 2D space into boxes with length  $n^{-\frac{1}{4}}$ , average in each box

**What's our estimation error?**

Informally, we have  $\sqrt{n}$  boxes, each box gets  $\sqrt{n}$  samples.

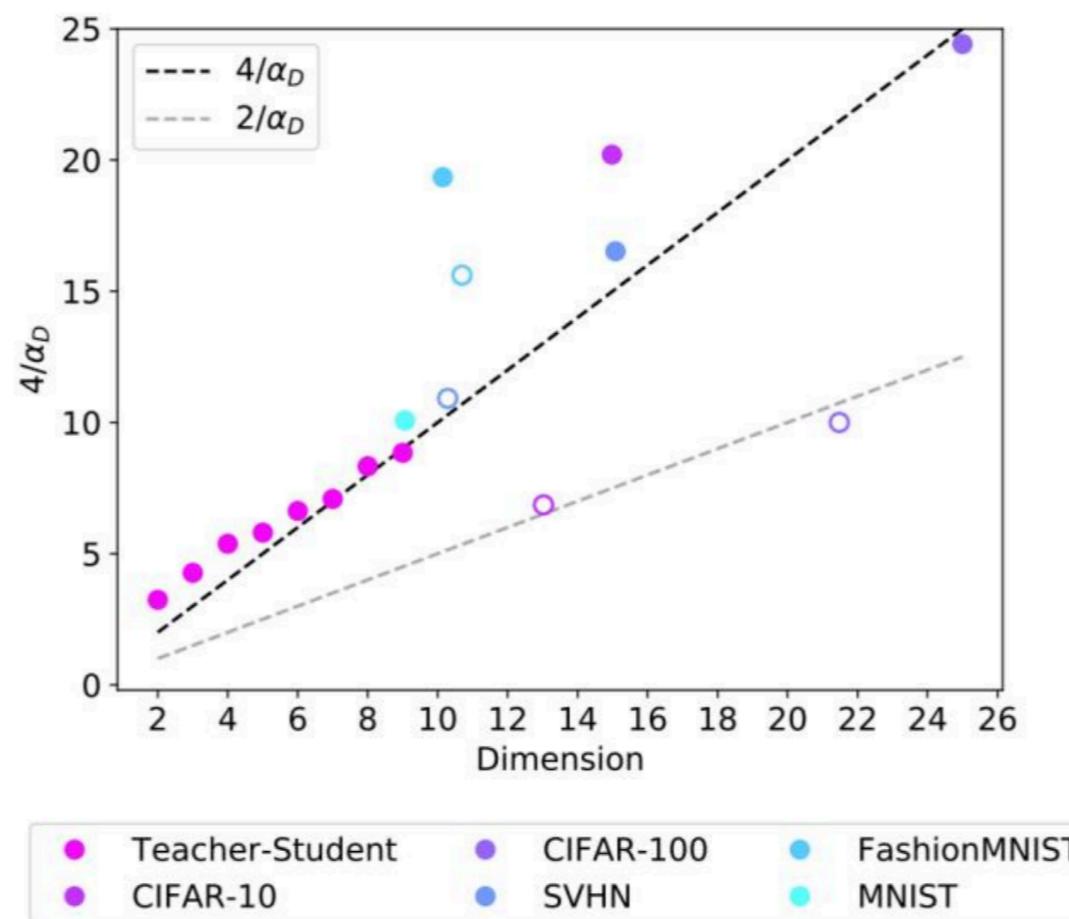
$$Error \approx \frac{1}{\sqrt{n}} + (\text{other smoothness terms})$$

In  $d$ -dimensions, this becomes  $Error = n^{-1/d}$  - **This means scaling is  $y = -\frac{1}{d}x + C$**

**Takeaway:** flexible ‘nonparametric’ learning has dimension dependent scaling laws.

# Intrinsic dimensionality theory of data scaling laws

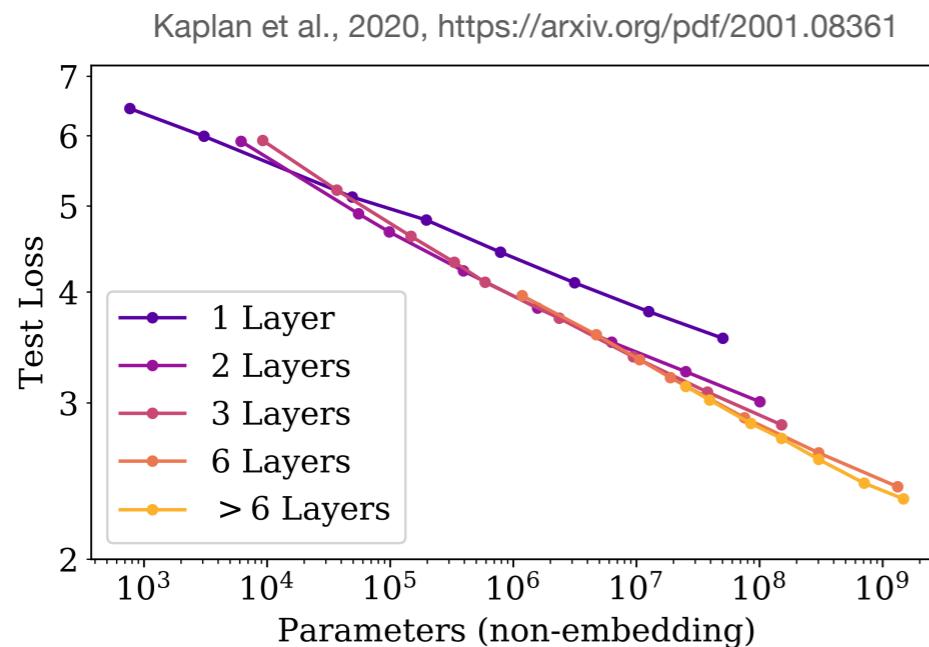
1. Scaling laws arise due to polynomial rates of learning  $\frac{1}{n^\alpha}$
2. The slope  $\alpha$  is closely connected to the *intrinsic dimensionality* of the data.



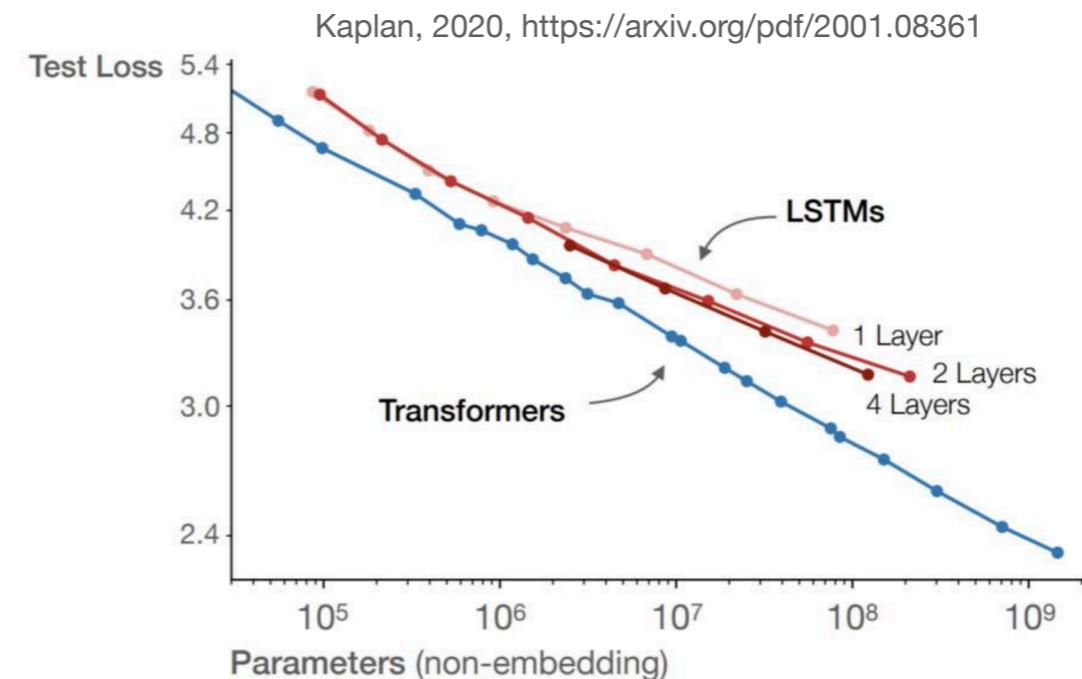
Some recent work (Bahri+ 2021) have tried to verify this empirically

# X-axis of Scaling Laws Can Be Different (# of parameters)

Best # of layers?

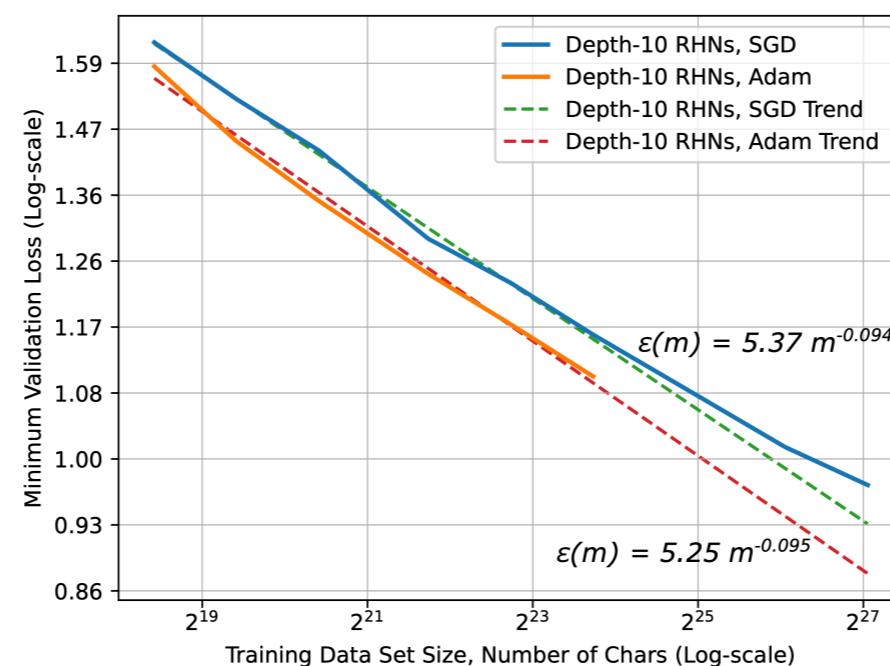


Best architecture?



Best optimizer?

Hestness, 2017, <https://arxiv.org/pdf/1712.00409>

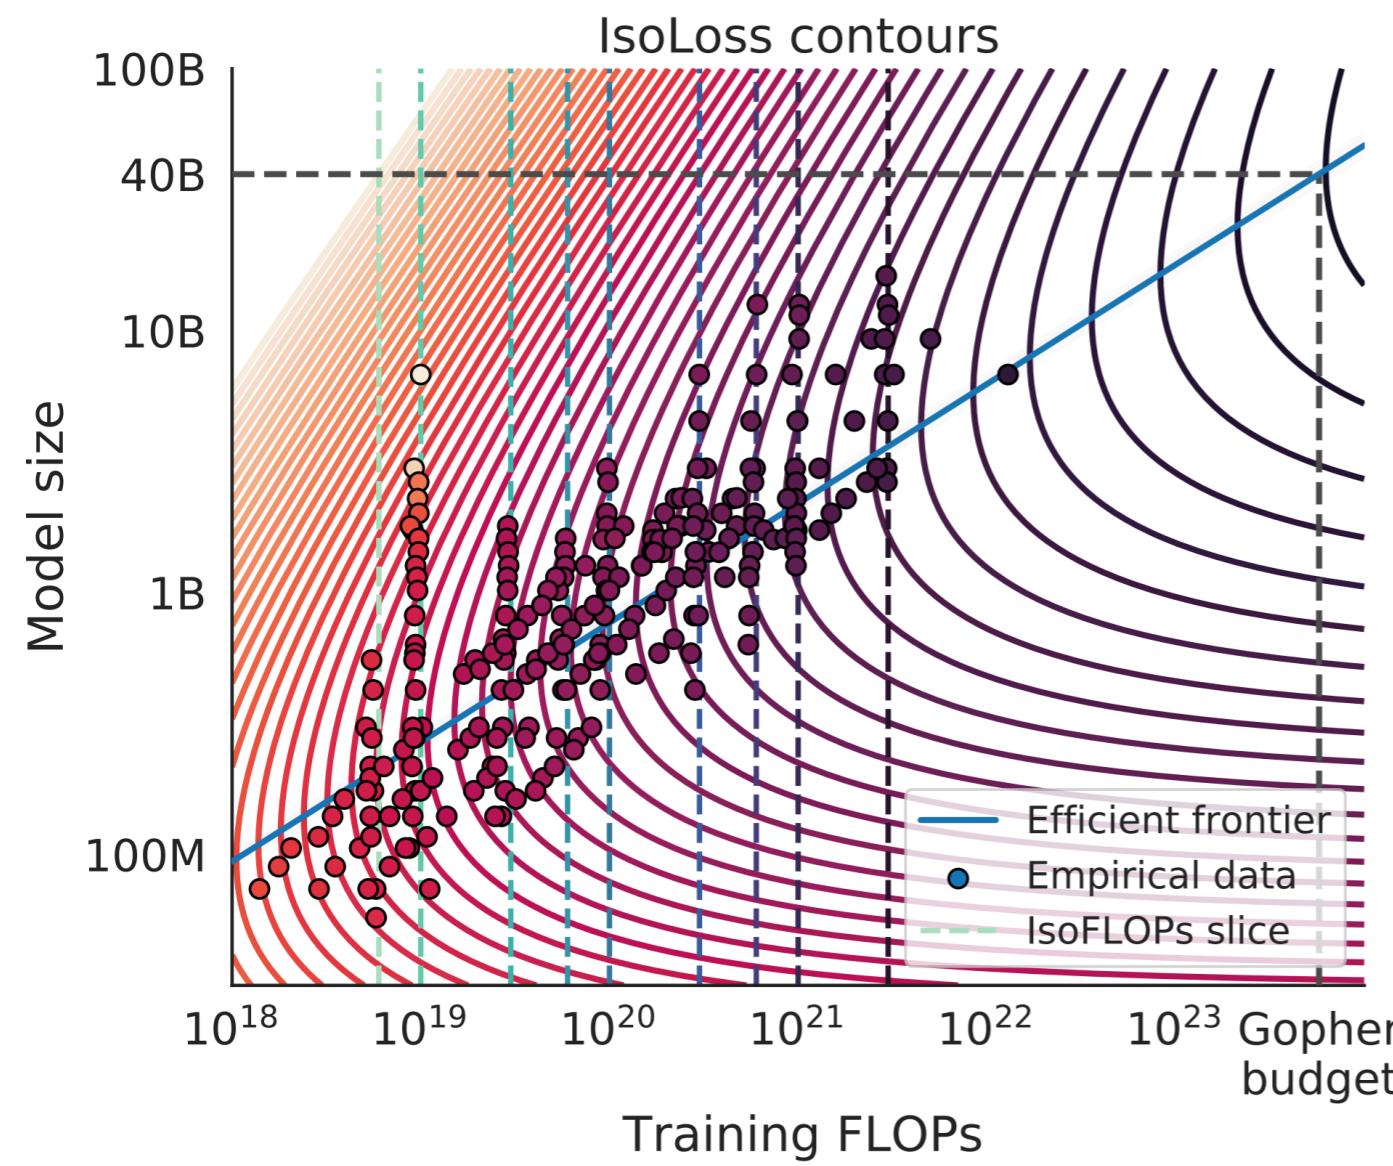


# Joint parameter-data scaling Law. How to scale?

$$\hat{L}(N, D) \triangleq E + \frac{A}{N^\alpha} + \frac{B}{D^\beta}$$

$N$  Number of Tokens  
 $D$  Number of Parameters

Hoffmann, 2022, <https://arxiv.org/pdf/2203.15556>

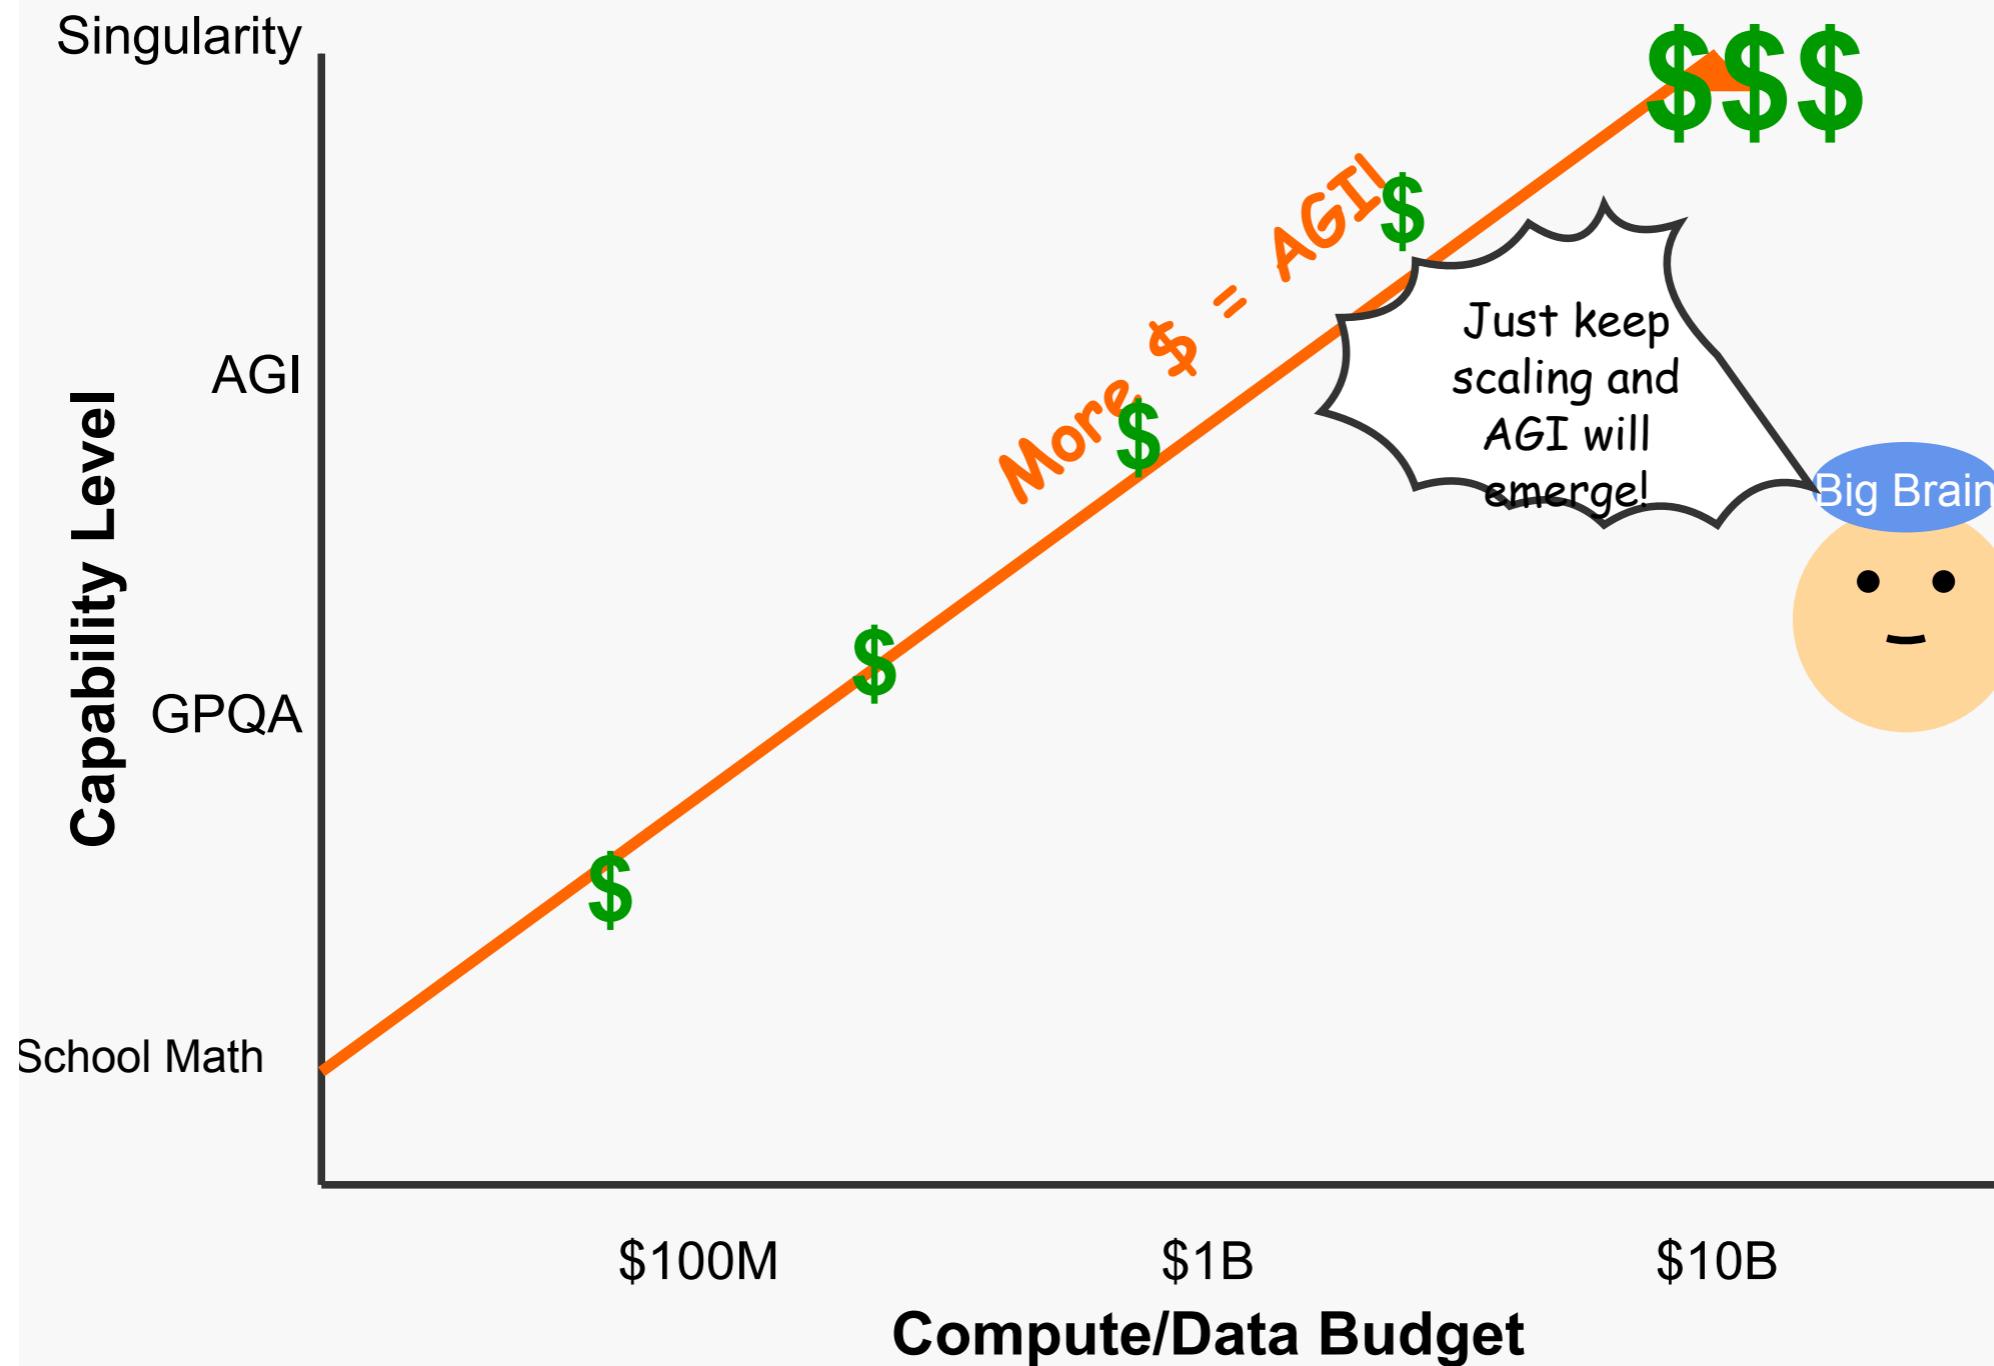


Empirically  $\alpha \approx \beta$

Scale data and model size **proportionally**

# AGI Soon?

## The "Scaling Solves Everything" Belief



## Putting on reviewer 2's hat:

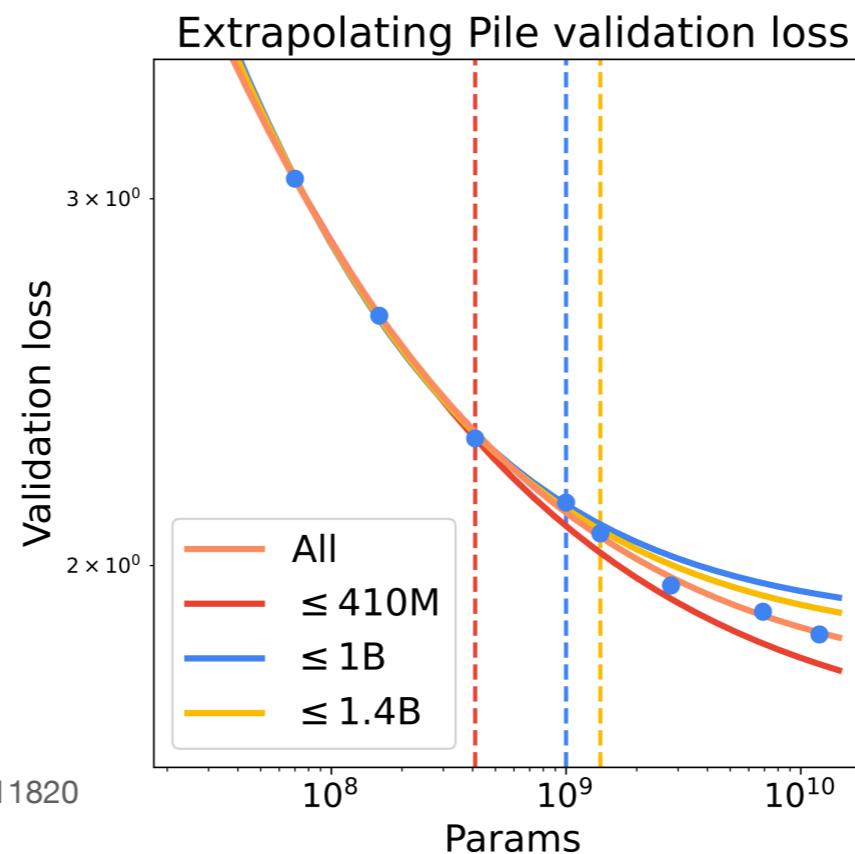
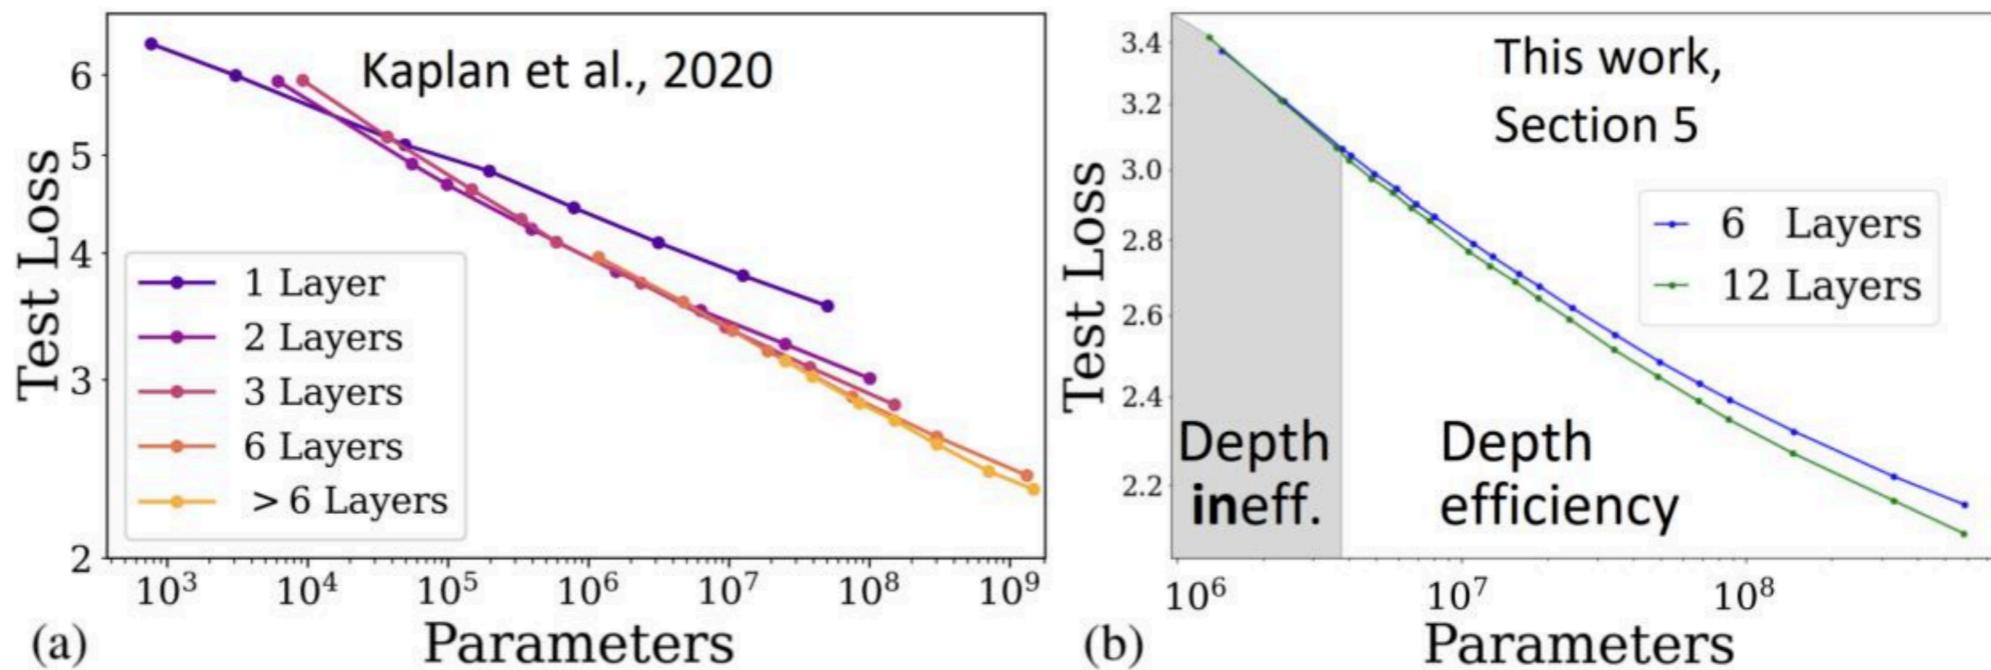
1. How well does it extrapolate? Does best parameter remain fixed at all scale?
2. What is x? Are all parameters/data equal?
3. What is y? What's the scaling behaviour on task accuracy? Or on OOD data
4. How do parameters of scaling law depend on architecture, or on the relationship between train/test data?

$$L = L_\epsilon + \beta n^{-\alpha}$$

$\alpha = f(\text{architecture, relevance of train to test data})?$

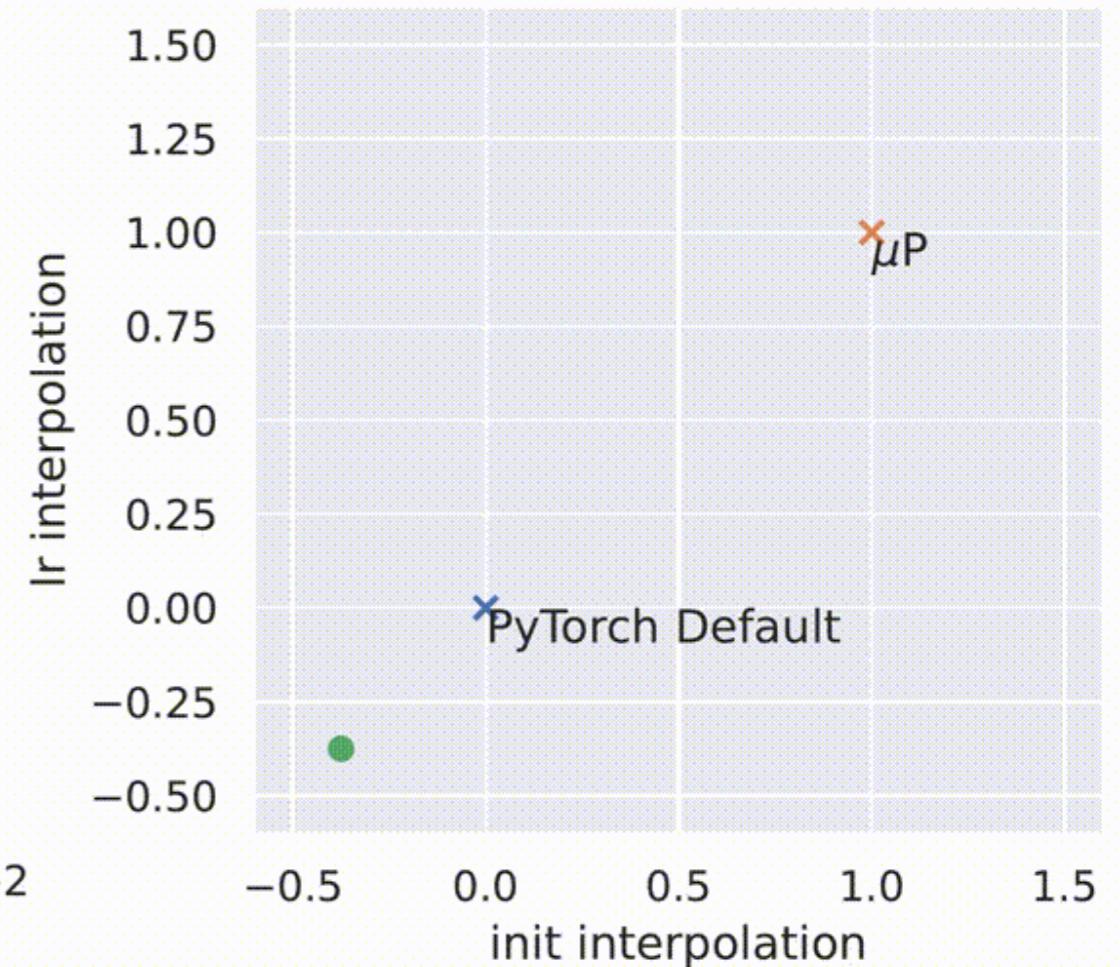
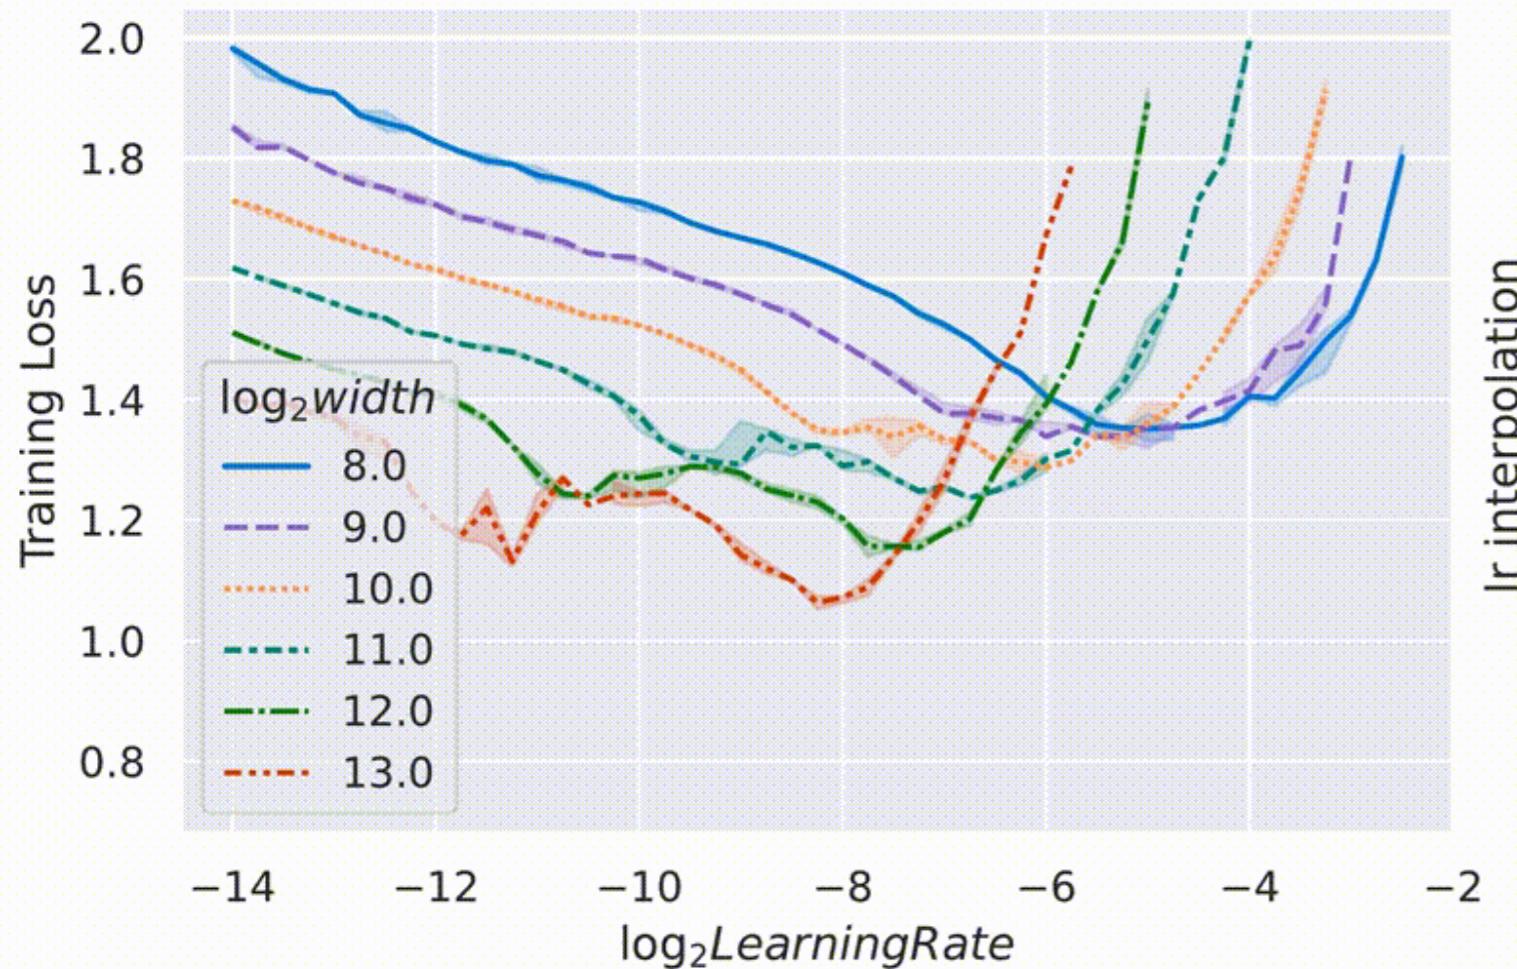
# Scaling Law's extrapolation can lead you astray

Levine et. al, 2021



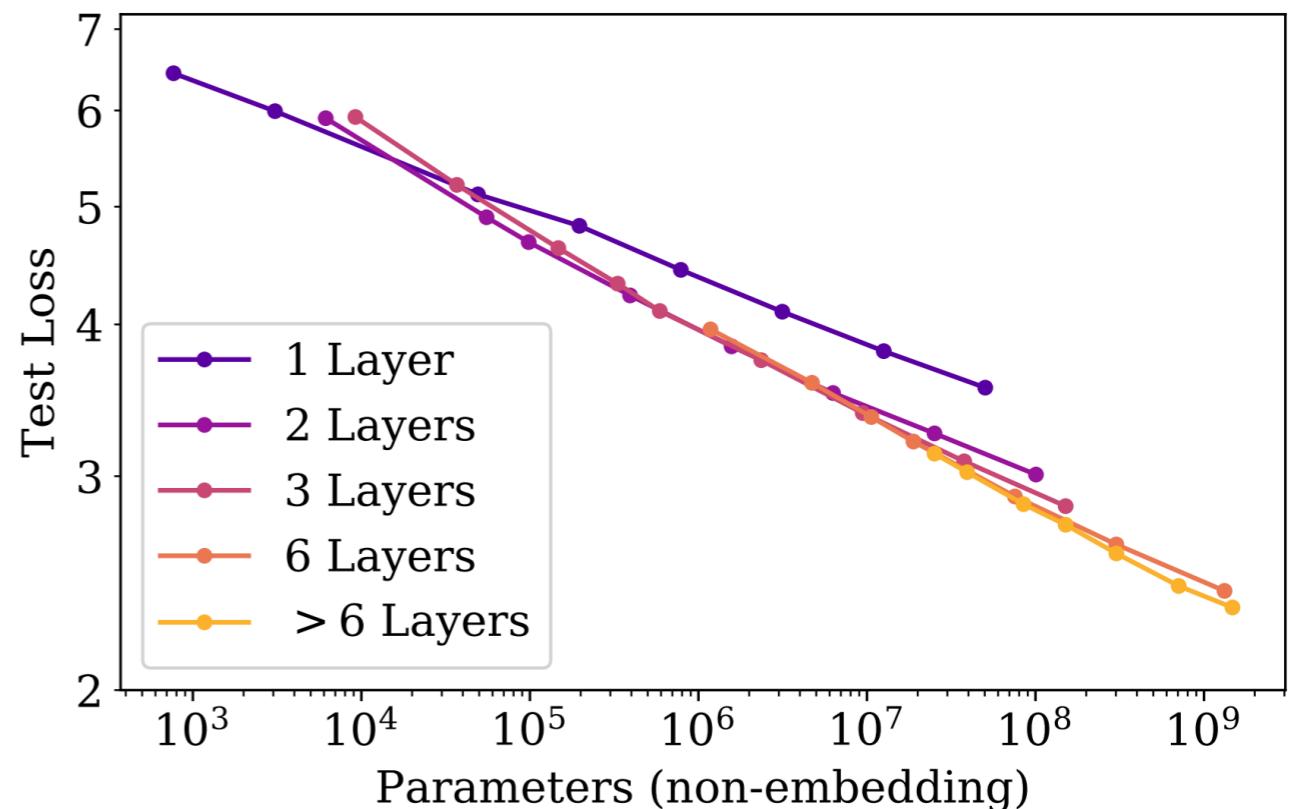
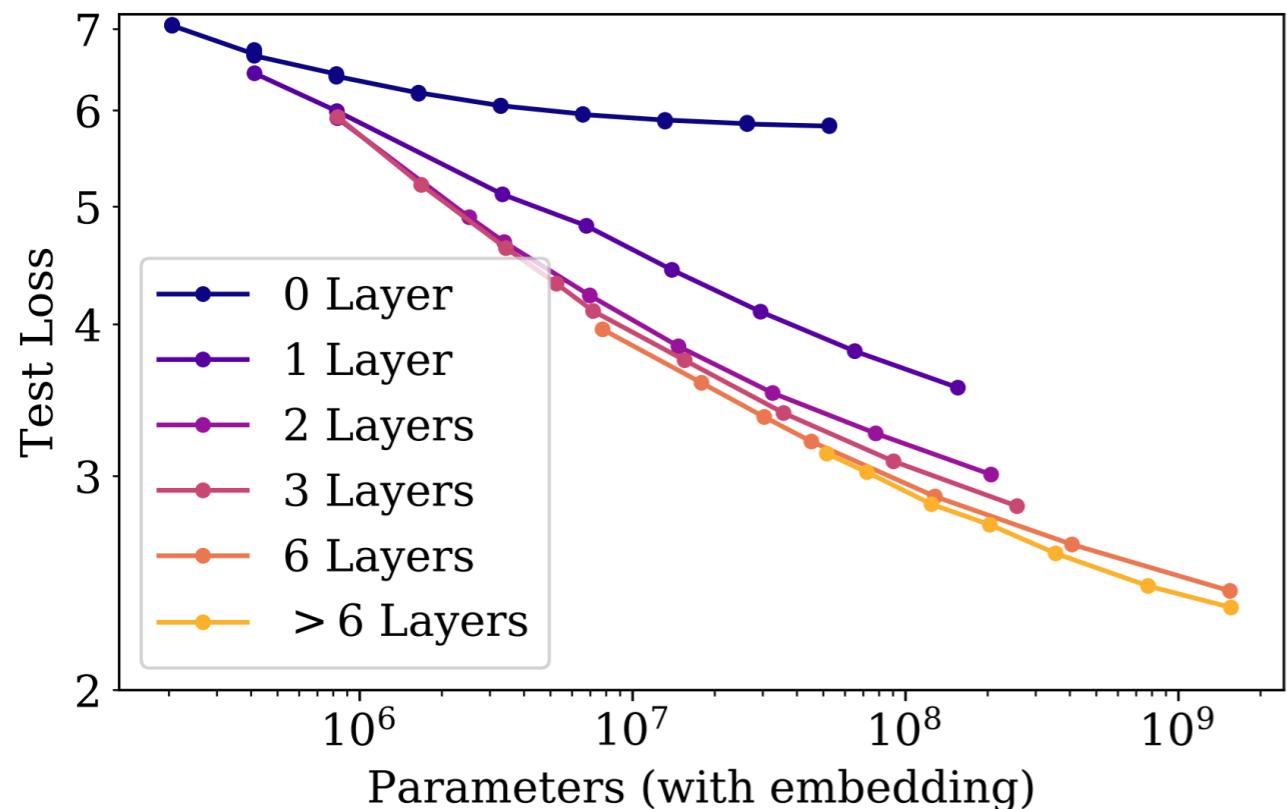
# Scaling Law's extrapolation can lead you astray

<https://www.microsoft.com/en-us/research/blog/μtransfer-a-technique-for-hyperparameter-tuning-of-enormous-neural-networks/>



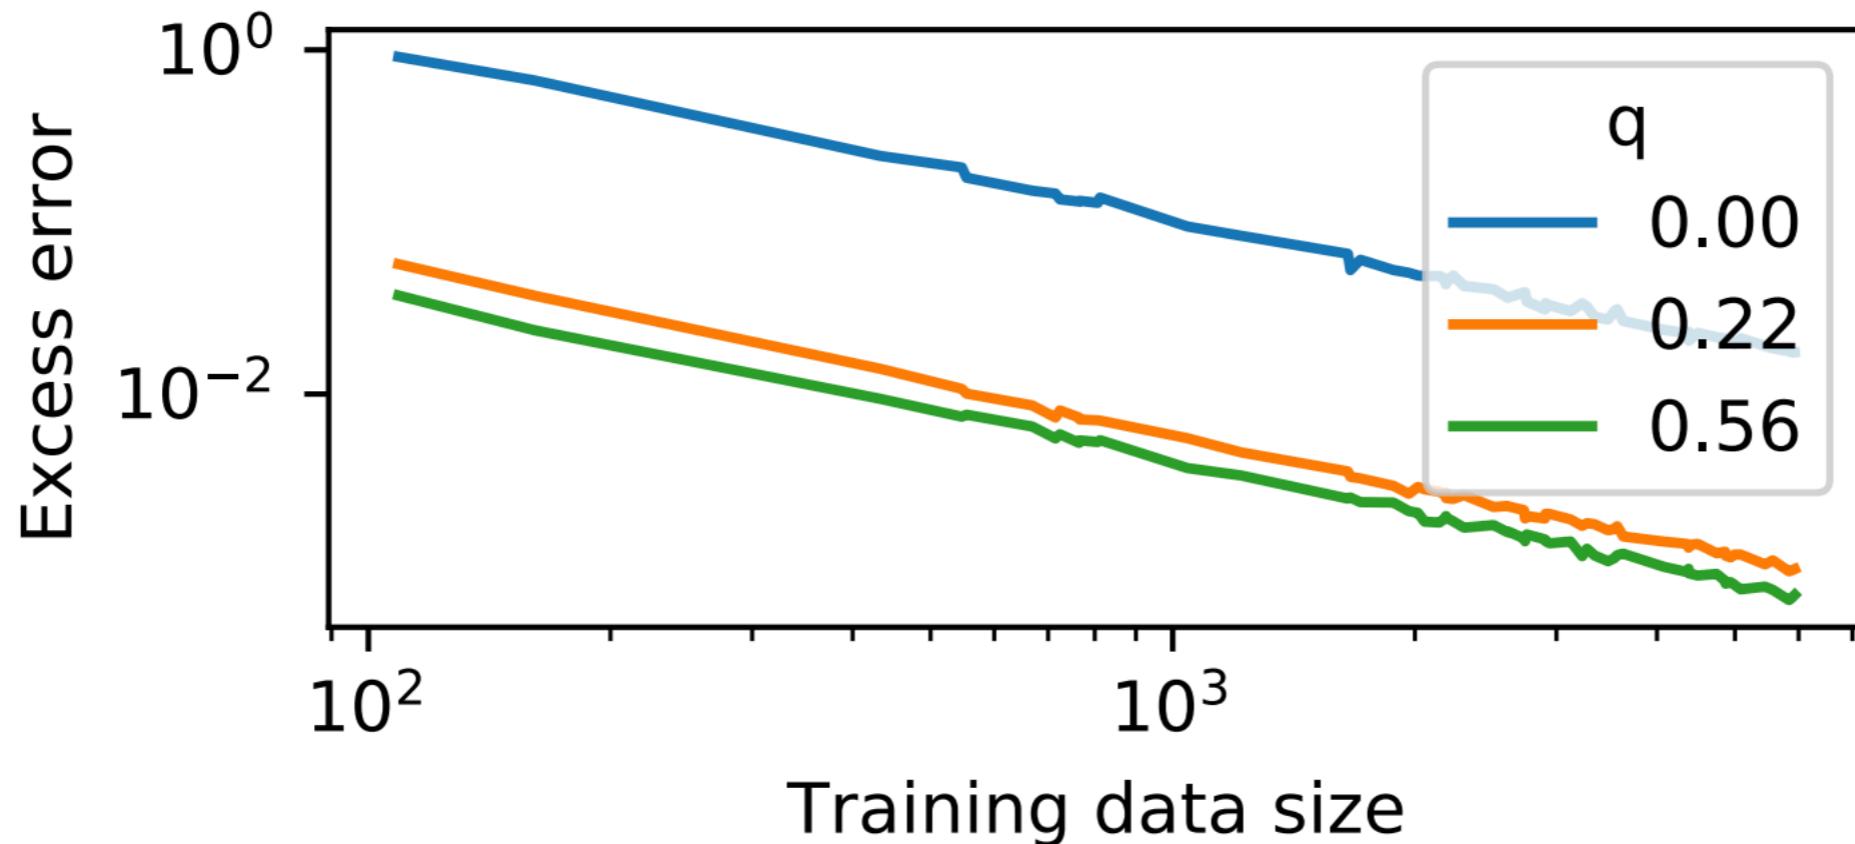
# Are all parameters equal?

Kaplan, 2020, <https://arxiv.org/pdf/2001.08361>



# Are all data equal? Intercept changes but exponent does not

Hashimoto, 2021, <https://proceedings.mlr.press/v139/hashimoto21a/hashimoto21a.pdf>



$q$ : proportion of one of the two training data

$$\log(L(n, q)) \approx \log(V(n, q)) := \alpha(q) \log(n) + C(q). \quad \text{X}$$

$$\log(L(n, q)) \approx \log(V(n, q)) := -\alpha \log(n) + \log(C(q)). \quad \checkmark$$

Data composition does not affect the slope?

# Are all data equal? Different data sources changes exponent

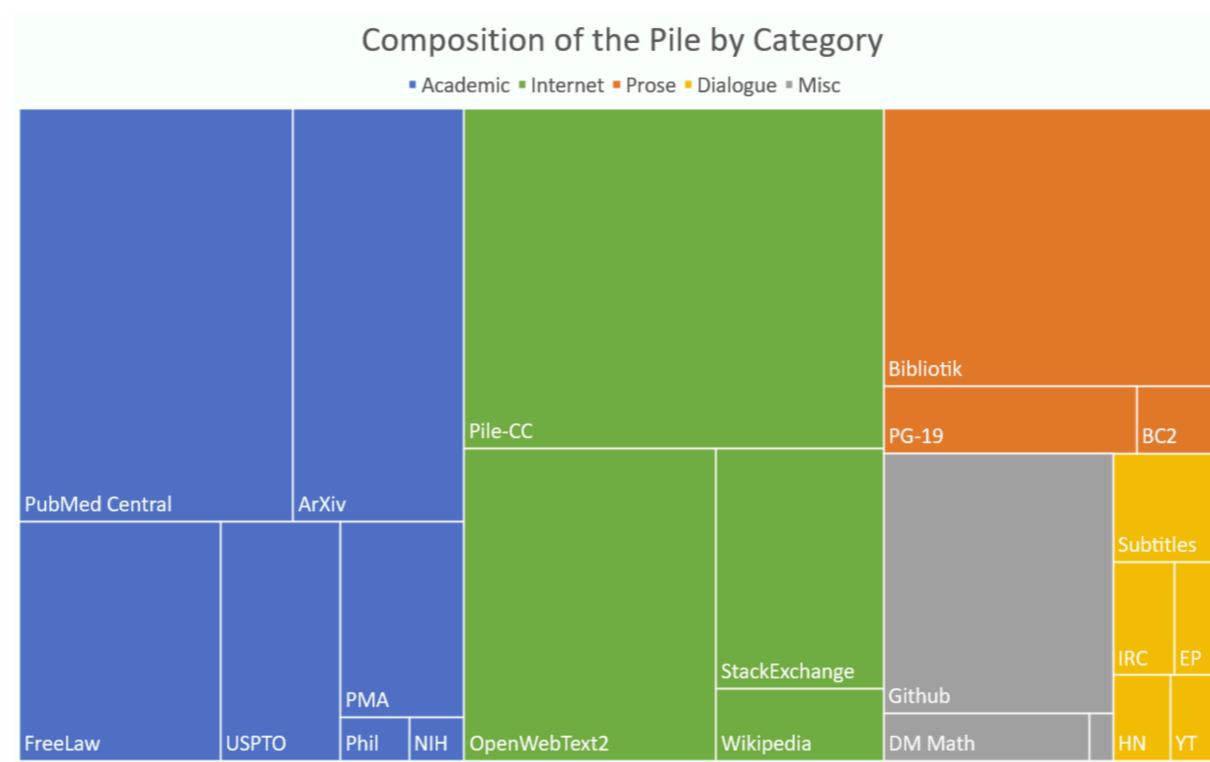


Figure 1: Treemap of Pile components by effective size.

$$L_i(r_{1\dots M}) = c_i + k_i \exp \left( \sum_{j=1}^M t_{ij} r_j \right)$$

i: Domain of validation data

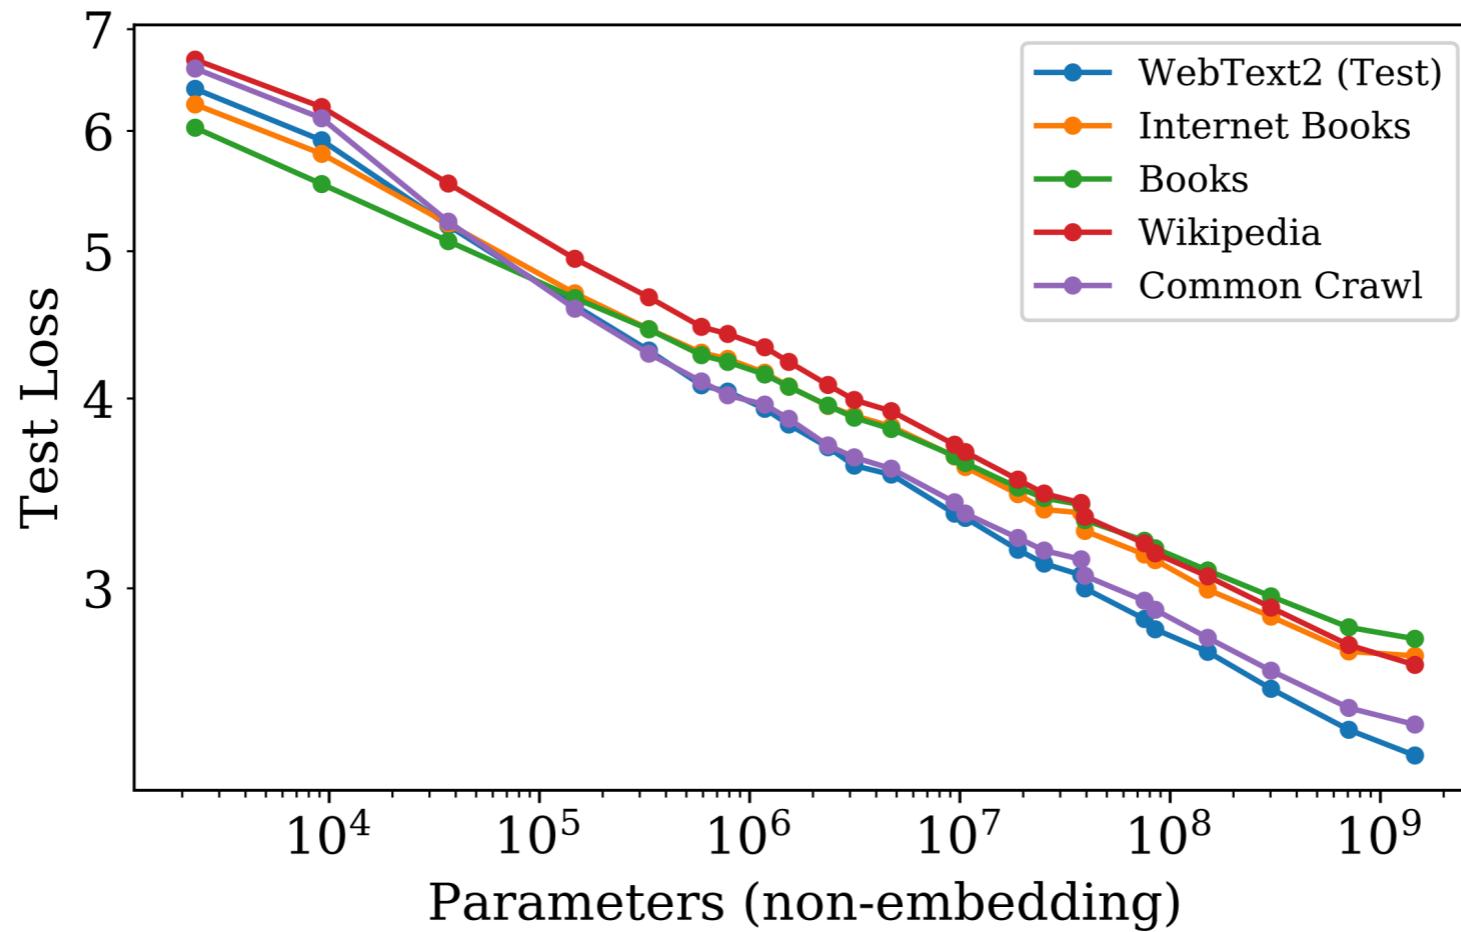
j: Domain of training data

r\_j: Proportion of training data from domain j

t\_ij: How much does training domain j helps validation domain I

# Are all data equal? Different target data changes exponent

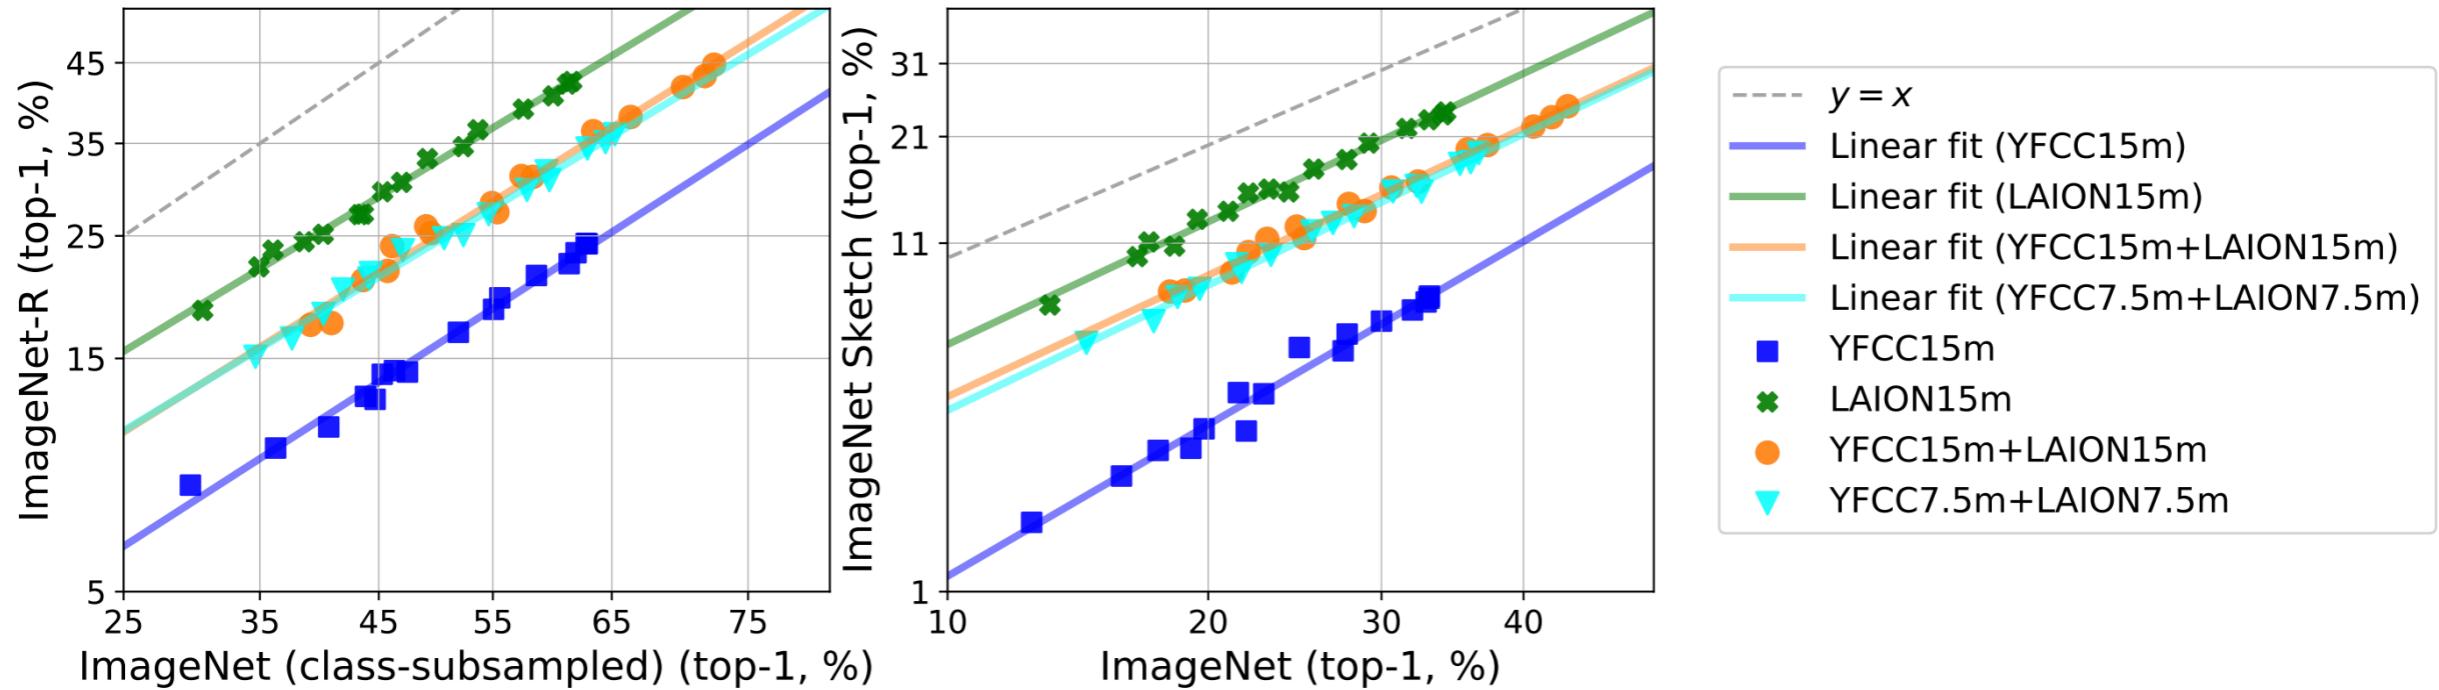
Kaplan et al., 2020, <https://arxiv.org/pdf/2001.08361>



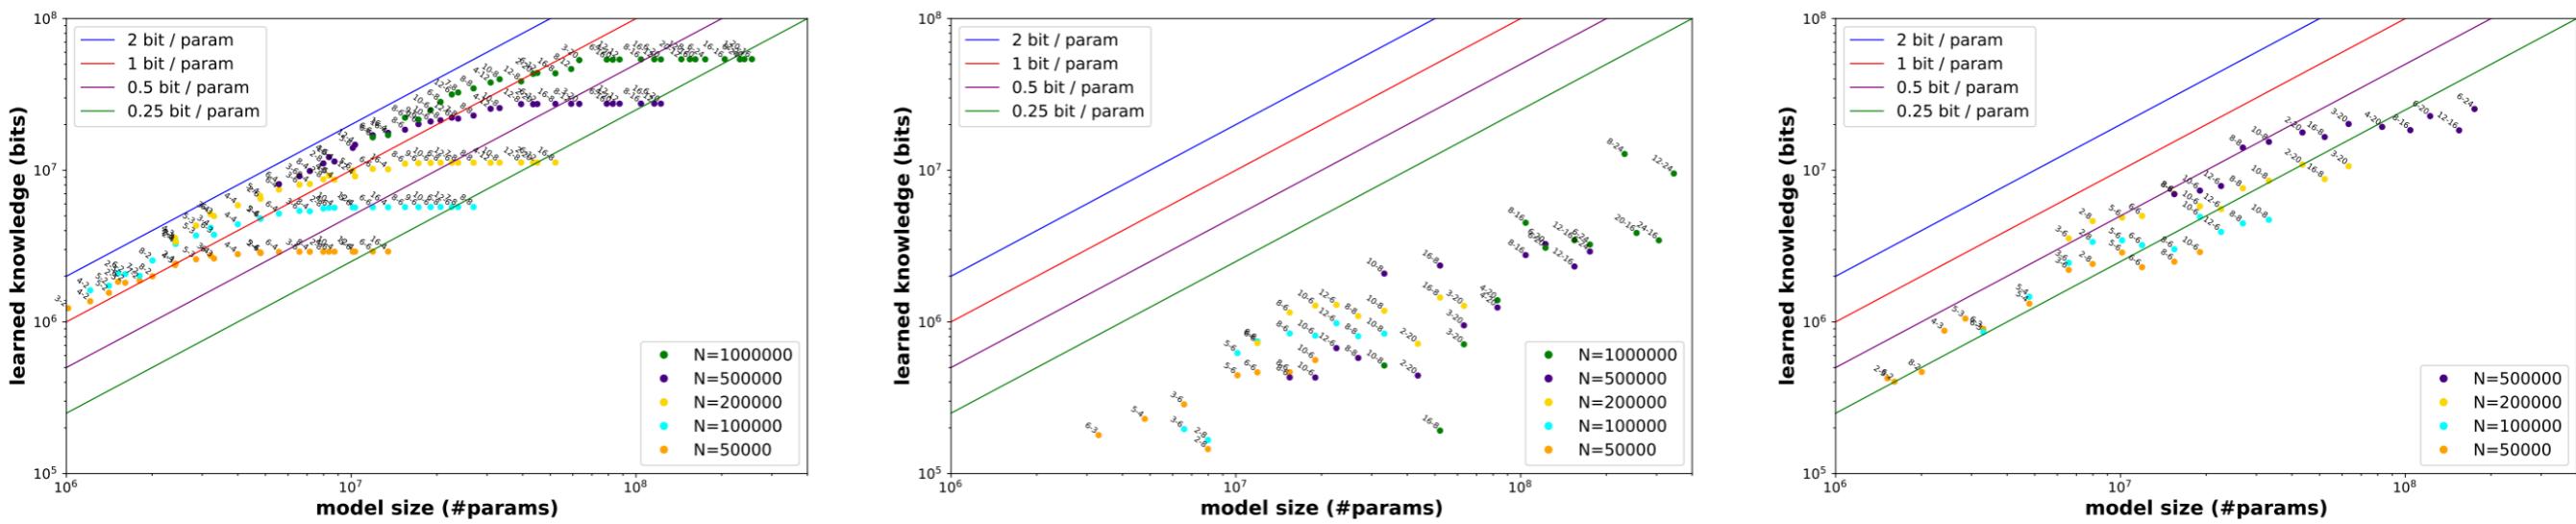
Clearly Webtext & Common Crawl have slopes different from Book's

# Are all data equal? Bad data can worsen your model

Nguyen et al., 2022, <https://arxiv.org/pdf/2208.05516>



Zhu et al., 2024, <https://arxiv.org/pdf/2404.05405>



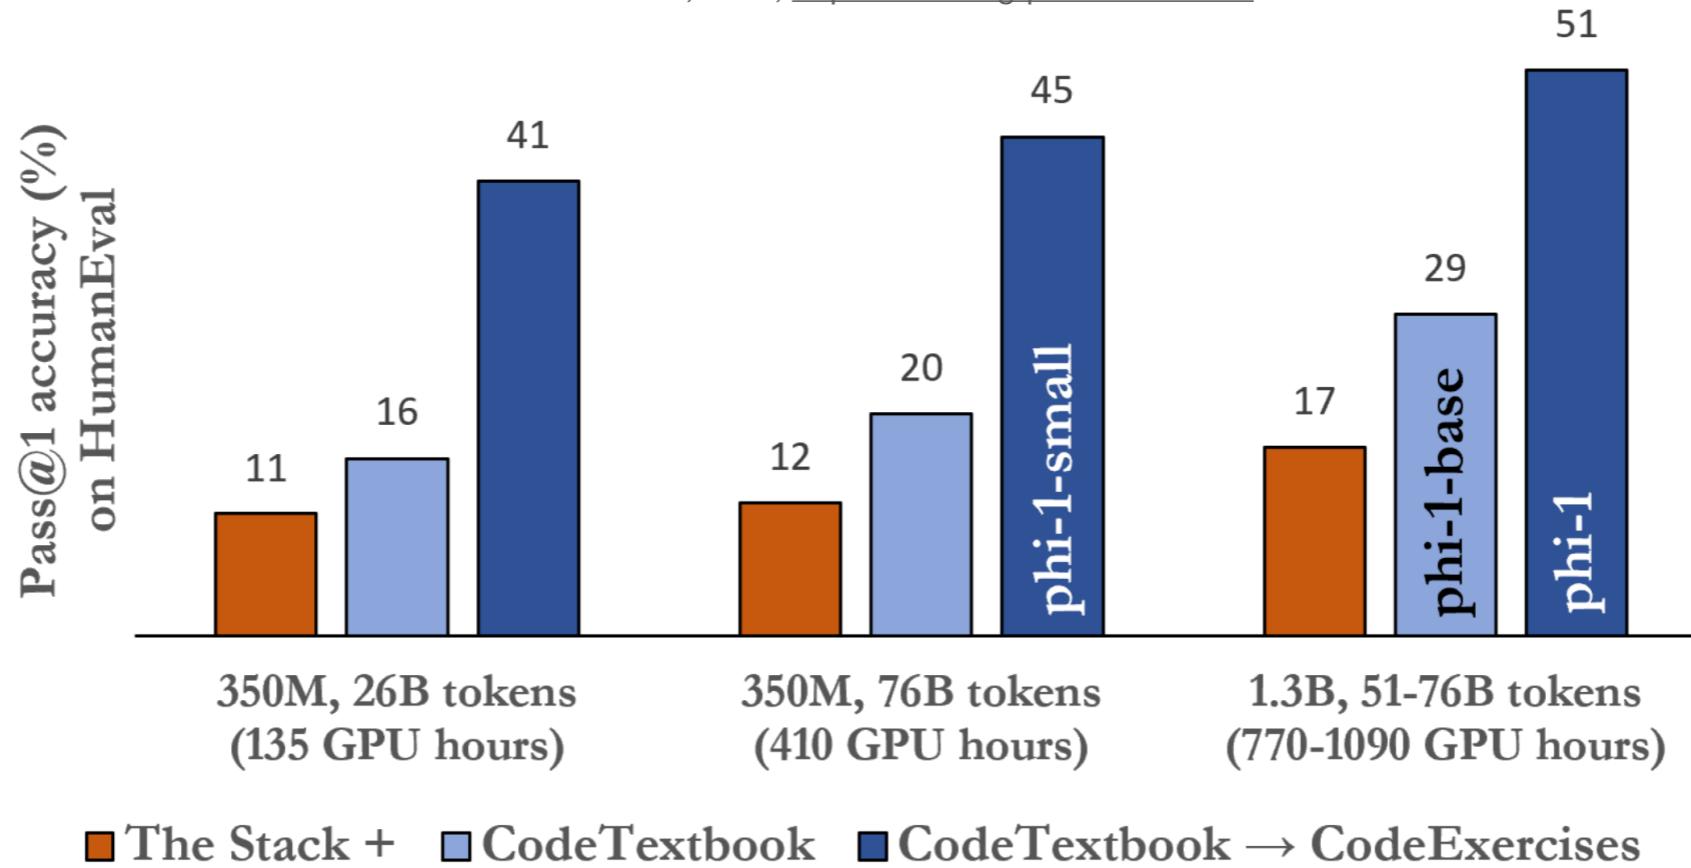
(a) no junk, 100 exposures

(b) 7/8 junk, 100 exposures

(c) 7/8 junk, 300 exposures

# Are all data equal? High-quality data gives you a lot more

Gunasekar et al., 2023, <https://arxiv.org/pdf/2306.11644>



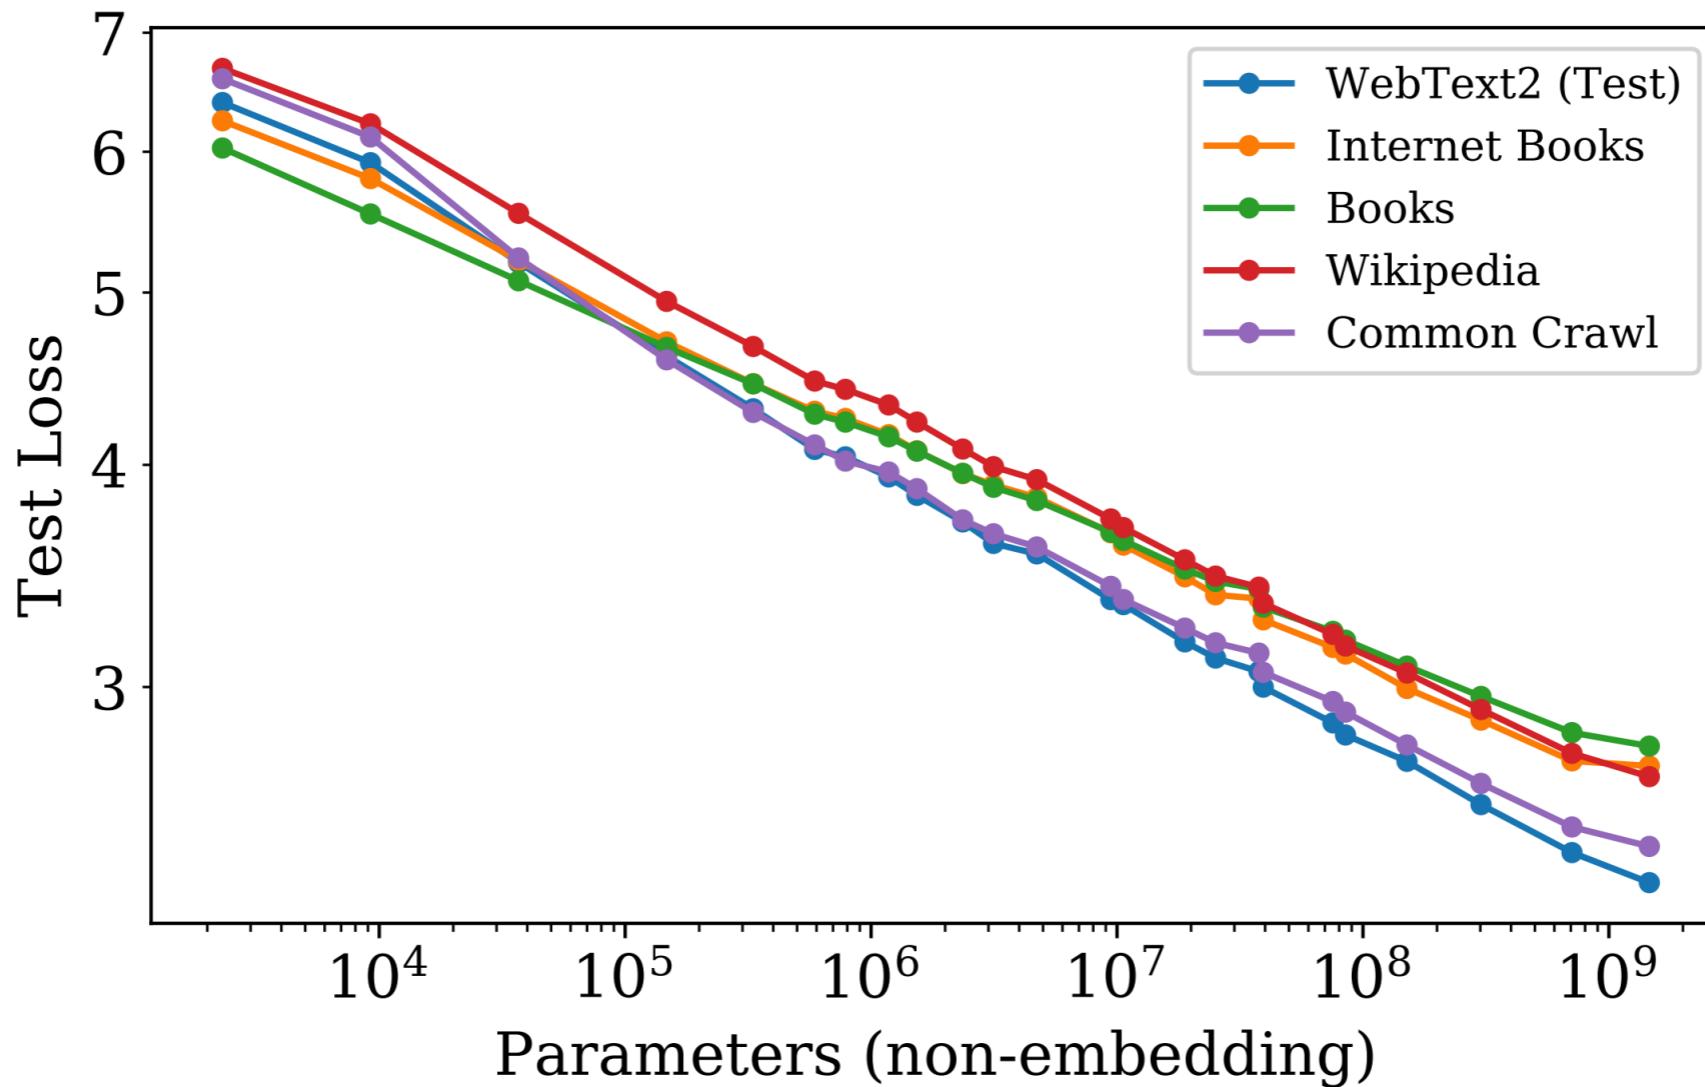
Tiny Stories: **10M**-sized model can generate coherent English  
when **125M** models (GPT-Neo, GPT-2) cannot.

Eldan et al., 2023, <https://arxiv.org/pdf/2305.07759>

# Does Scaling Law work out of distribution?

Trained on WebText. Evaluate on the rest.

Kaplan, 2020, <https://arxiv.org/pdf/2001.08361>

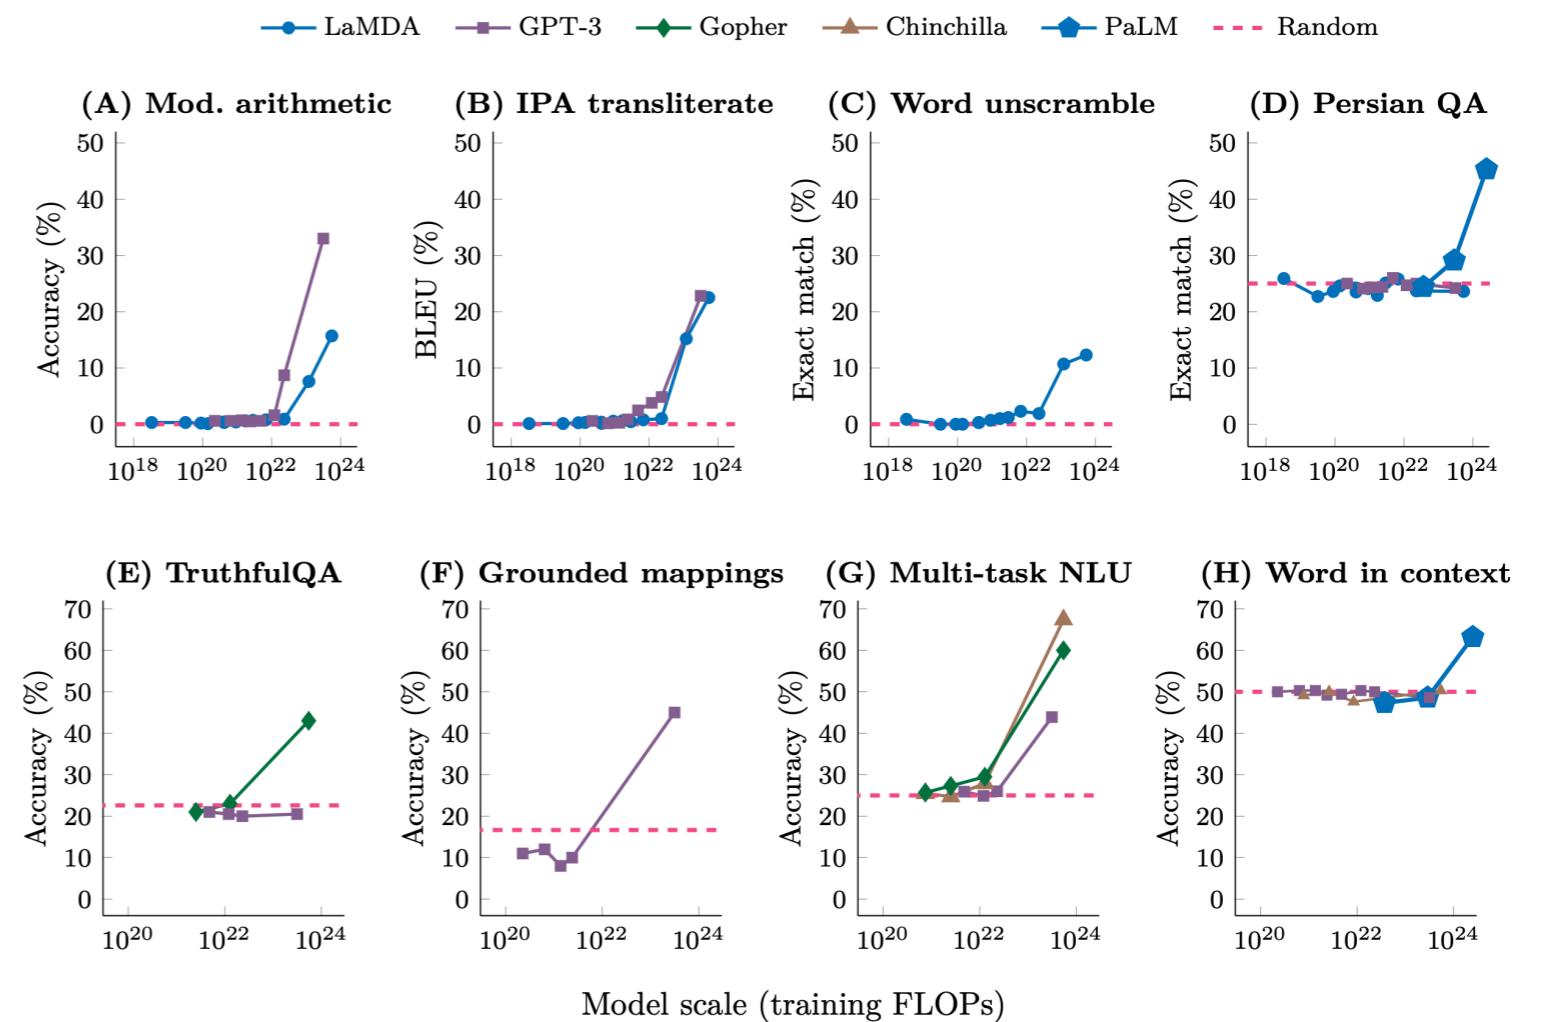


Scaling Law still holds, albeit with different intercept (and slope?)

# Does Scaling Law work for downstream tasks?

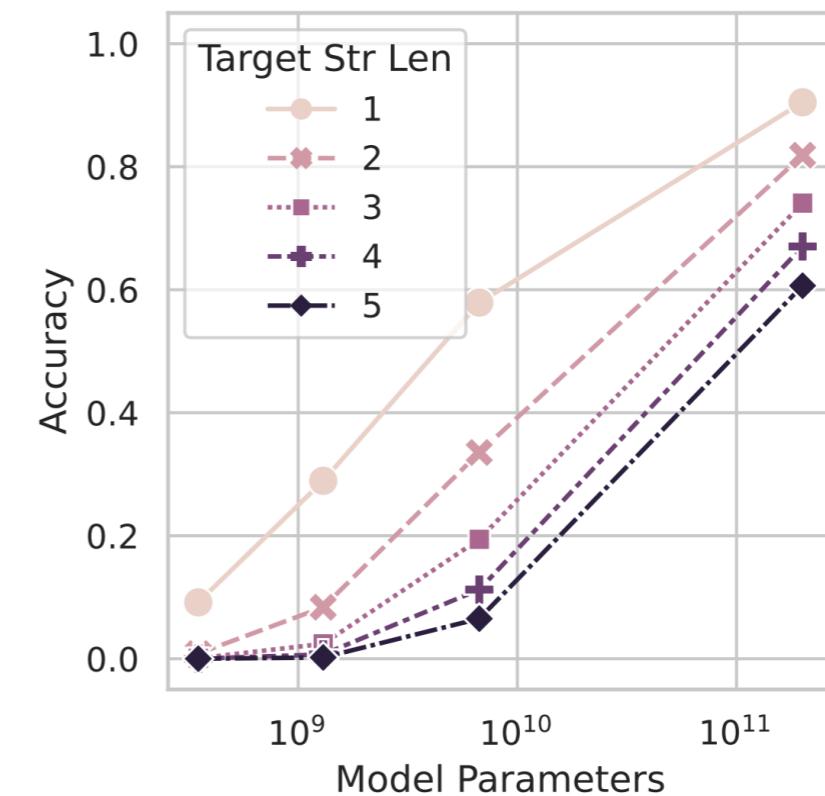
No, at first glance...

## Emergent Behaviour

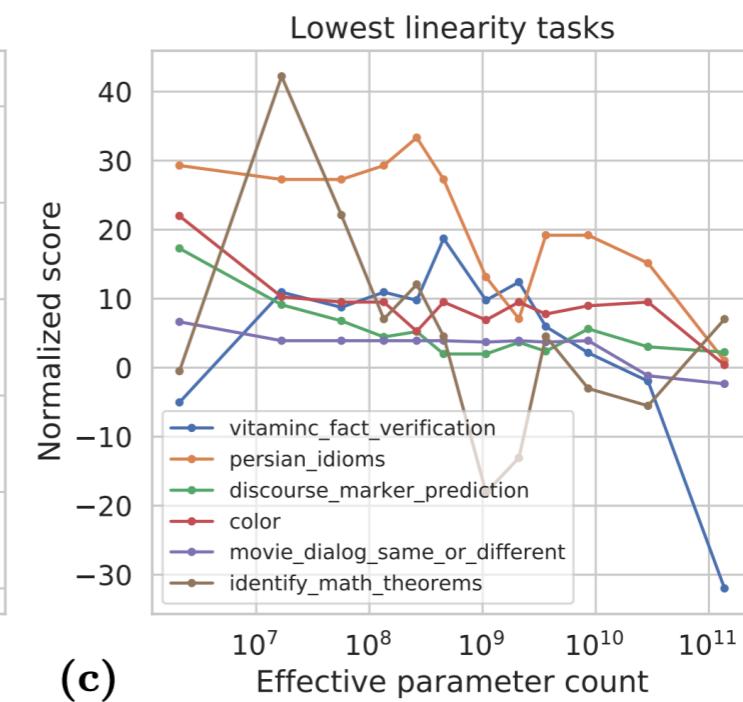
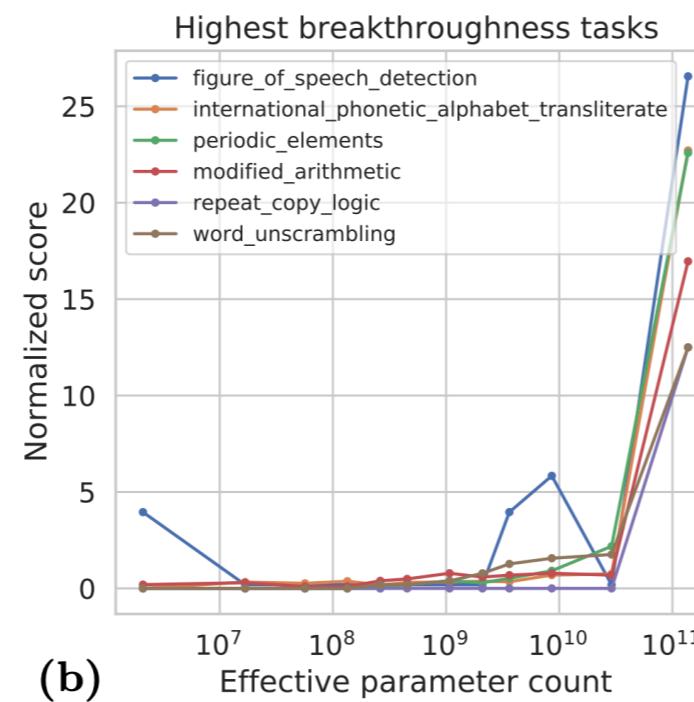
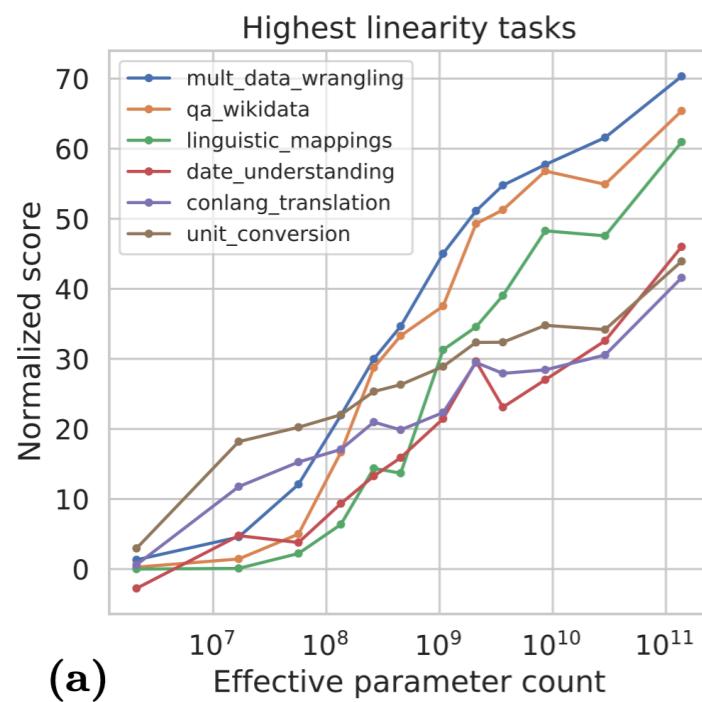


# Does Scaling Law work for downstream tasks?

Just because LLM needs to be correct multiple times



Task dependent:



# Recap on Scaling Law

- Surprisingly robust pattern. Has theoretical foundation.
- Doesn't always work. Need to carefully think about the axis.

# Motivating problem: hyperparameter costs

How can we solve this?

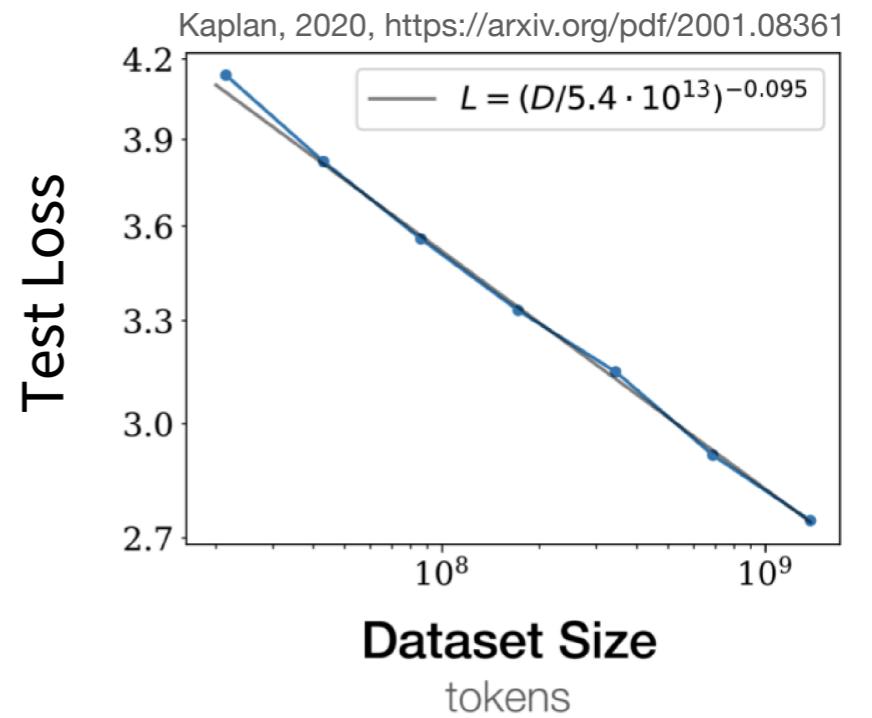
- Guess
- Grid Search
- Do small-scale experiments. Then “extrapolate”

## 1. Draw a line: **Scaling Law**

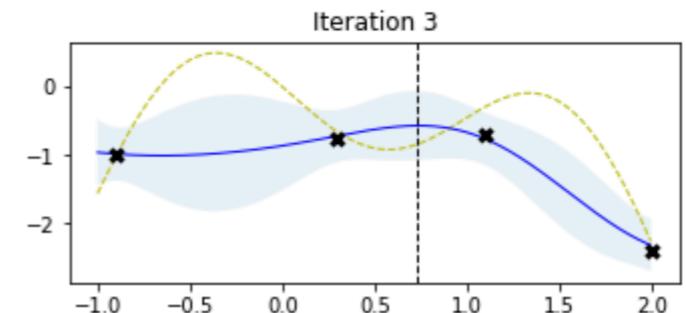
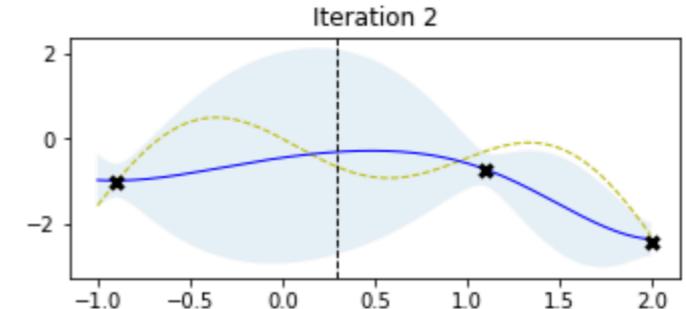
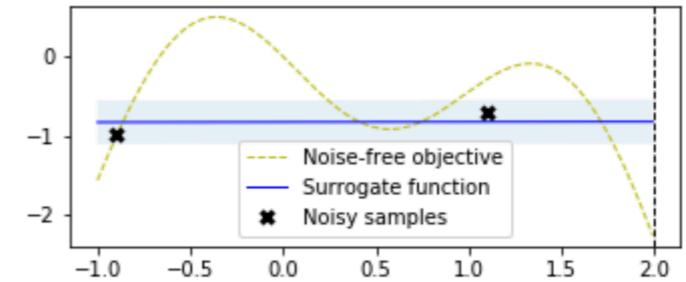
<https://stanford-cs324.github.io/winter2022/assets/pdfs/Scaling%20laws%20pdf.pdf>

## 2. (Multi-fidelity) Bayesian Optimization

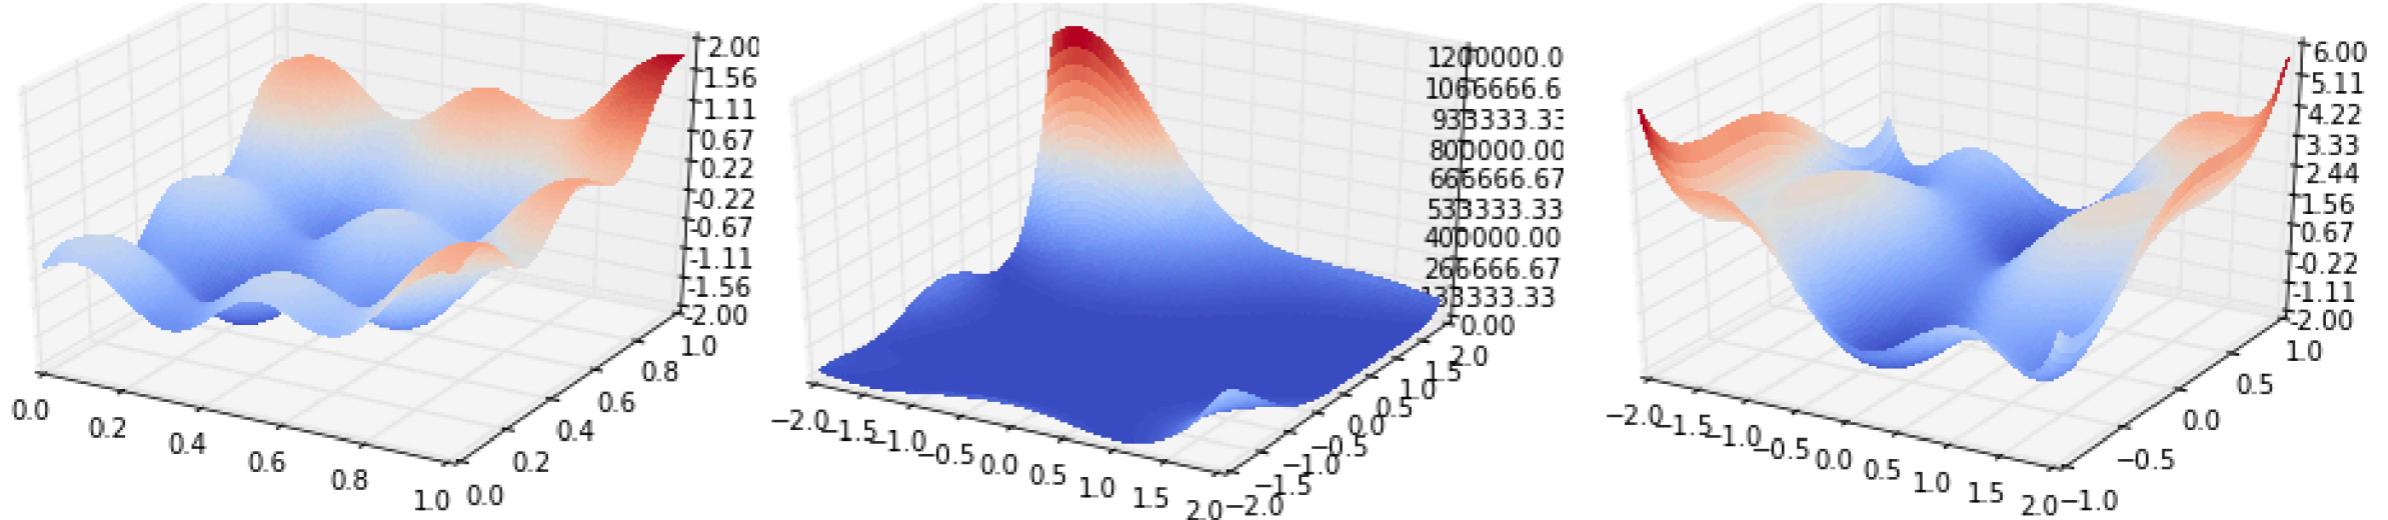
## 3. Update hyperparameters (specifically data) online



<https://krasserm.github.io/2018/03/21/bayesian-optimization/>



# Hyperparameter Optimization



$$x^* = \arg \min f(x)$$

- $f$  is unknown (performance of data)
- $x$  is hyperparameter (data mixture, optimizer, learning rate etc.)
- No gradients
- Evaluation of  $f$  is expensive

# Bayesian Optimization - Bayesian Statistics

$$P(\theta|D) = \frac{P(D|\theta) P(\theta)}{P(D)}$$

Diagram illustrating the Bayesian formula:

- Likelihood** (blue) is represented by  $P(D|\theta)$ .
- Prior** (green) is represented by  $P(\theta)$ .
- The product of Likelihood and Prior is divided by **Evidence** (black)  $P(D)$  to yield the **Posterior** (red)  $P(\theta|D)$ .

```
graph TD; L[P(D|θ)] --> P[P(θ)]; E[P(D)] --> P; P --> Pd[P(θ|D)]; L --- Pd; P --- Pd; E --- Pd;
```

$D$  data       $\theta$  something we do not observe

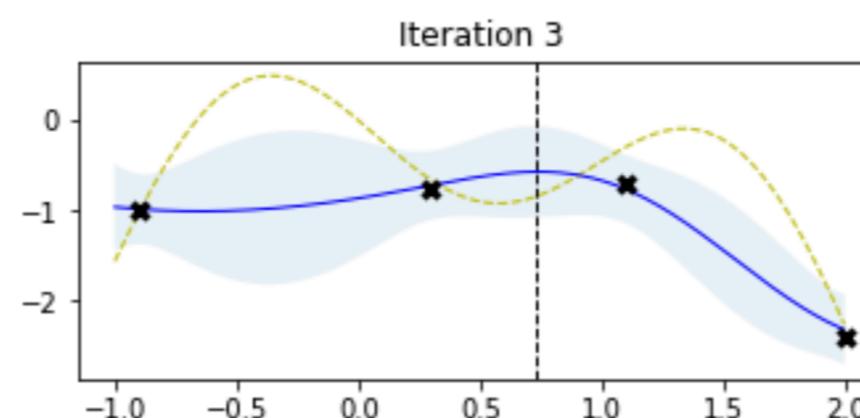
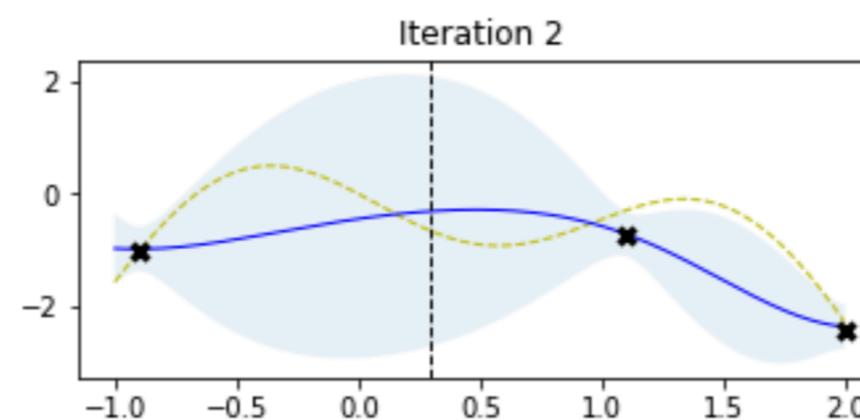
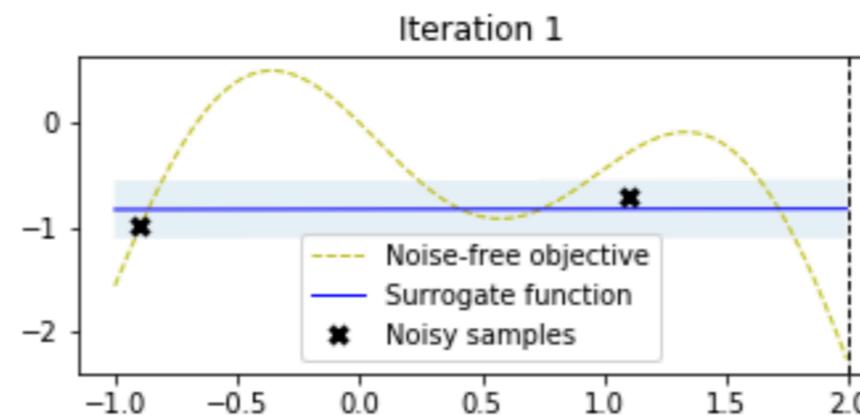
$P(\theta)$  initial belief of distribution of what we don't know

$P(D|\theta)$  the data generative process

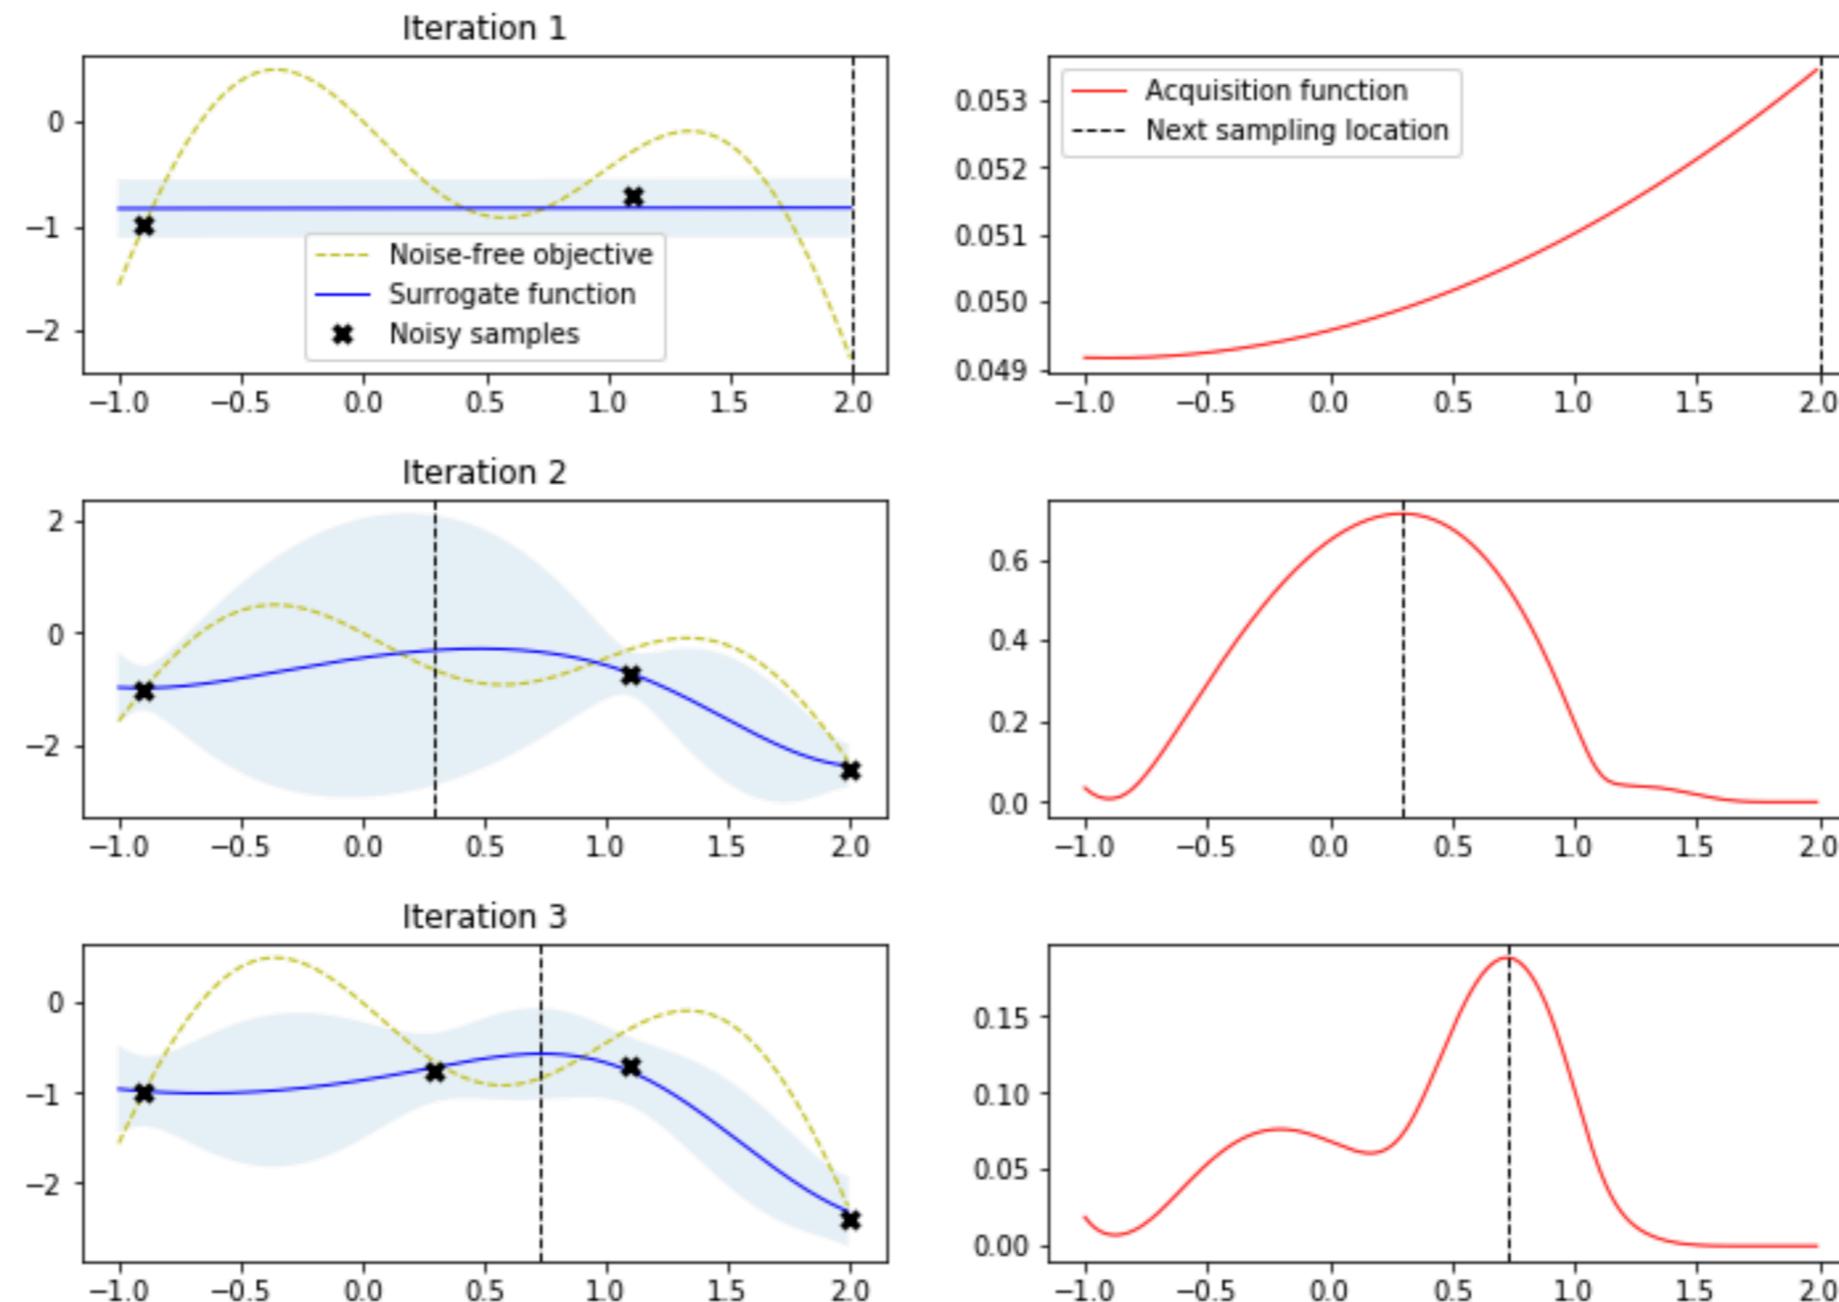
# Bayesian Optimization - Distribution of Function Given Data

$D$  data       $\theta$  functions

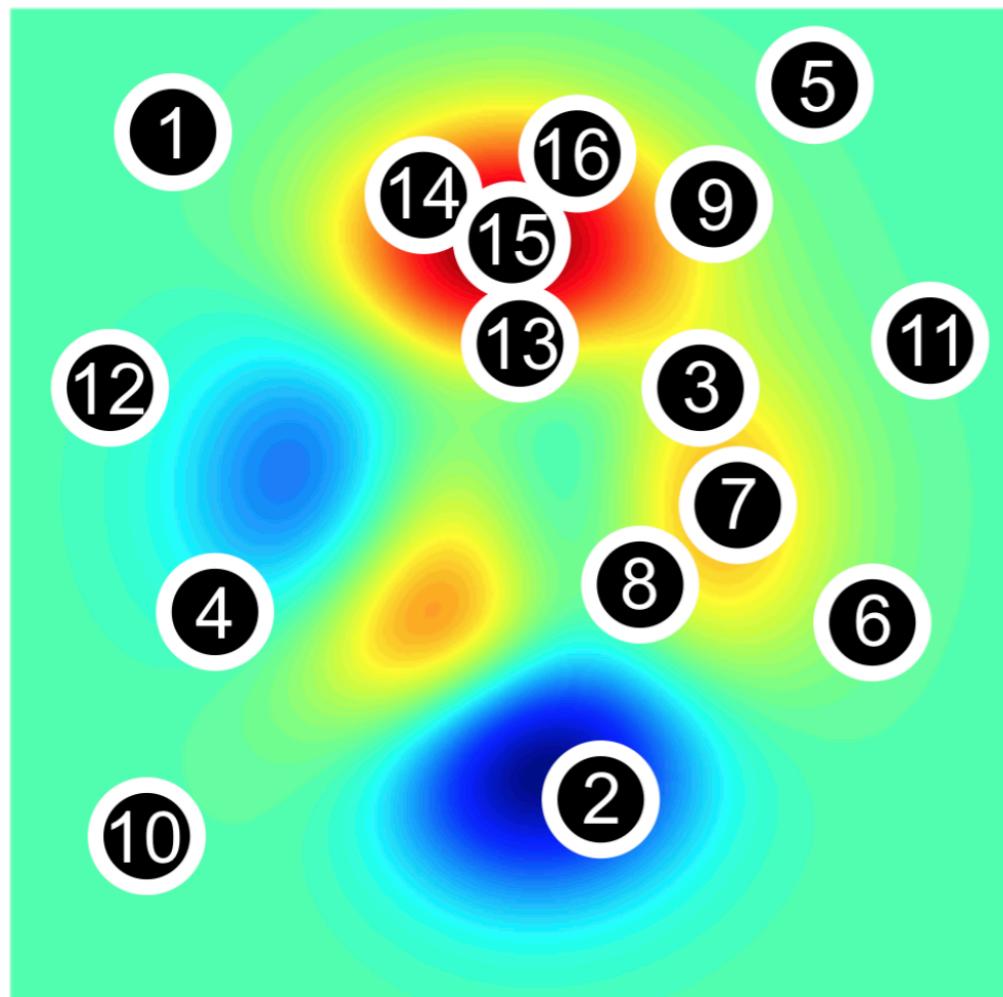
$P(\theta)$   $P(D|\theta)$  determined by [Gaussian process](#)



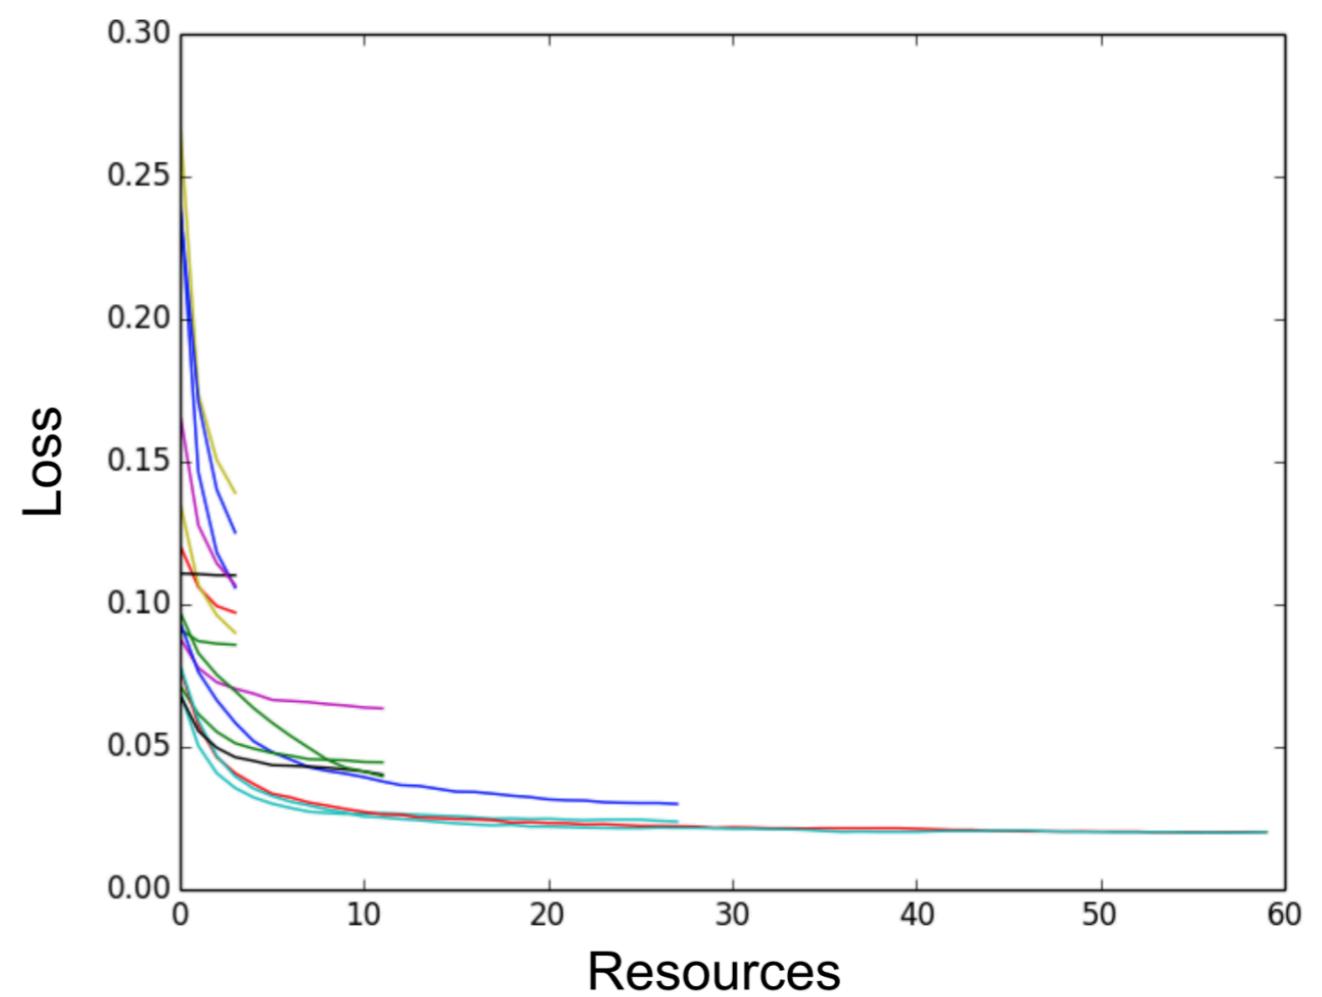
# Bayesian Optimization - What's the next point I should choose?



# Separate Idea: Multi-fidelity



(a) Configuration Selection



(b) Configuration Evaluation

# Multi-fidelity Bayesian Optimization

$$x^* = \arg \min f(x, s^*)$$

e.g. I want to train the best model for  $s^{*}=100$  steps

I can train any  $x$  with  $s < s^*$  steps with cost  $c(s)$

# Multi-fidelity Multi-scale Bayesian Optimization

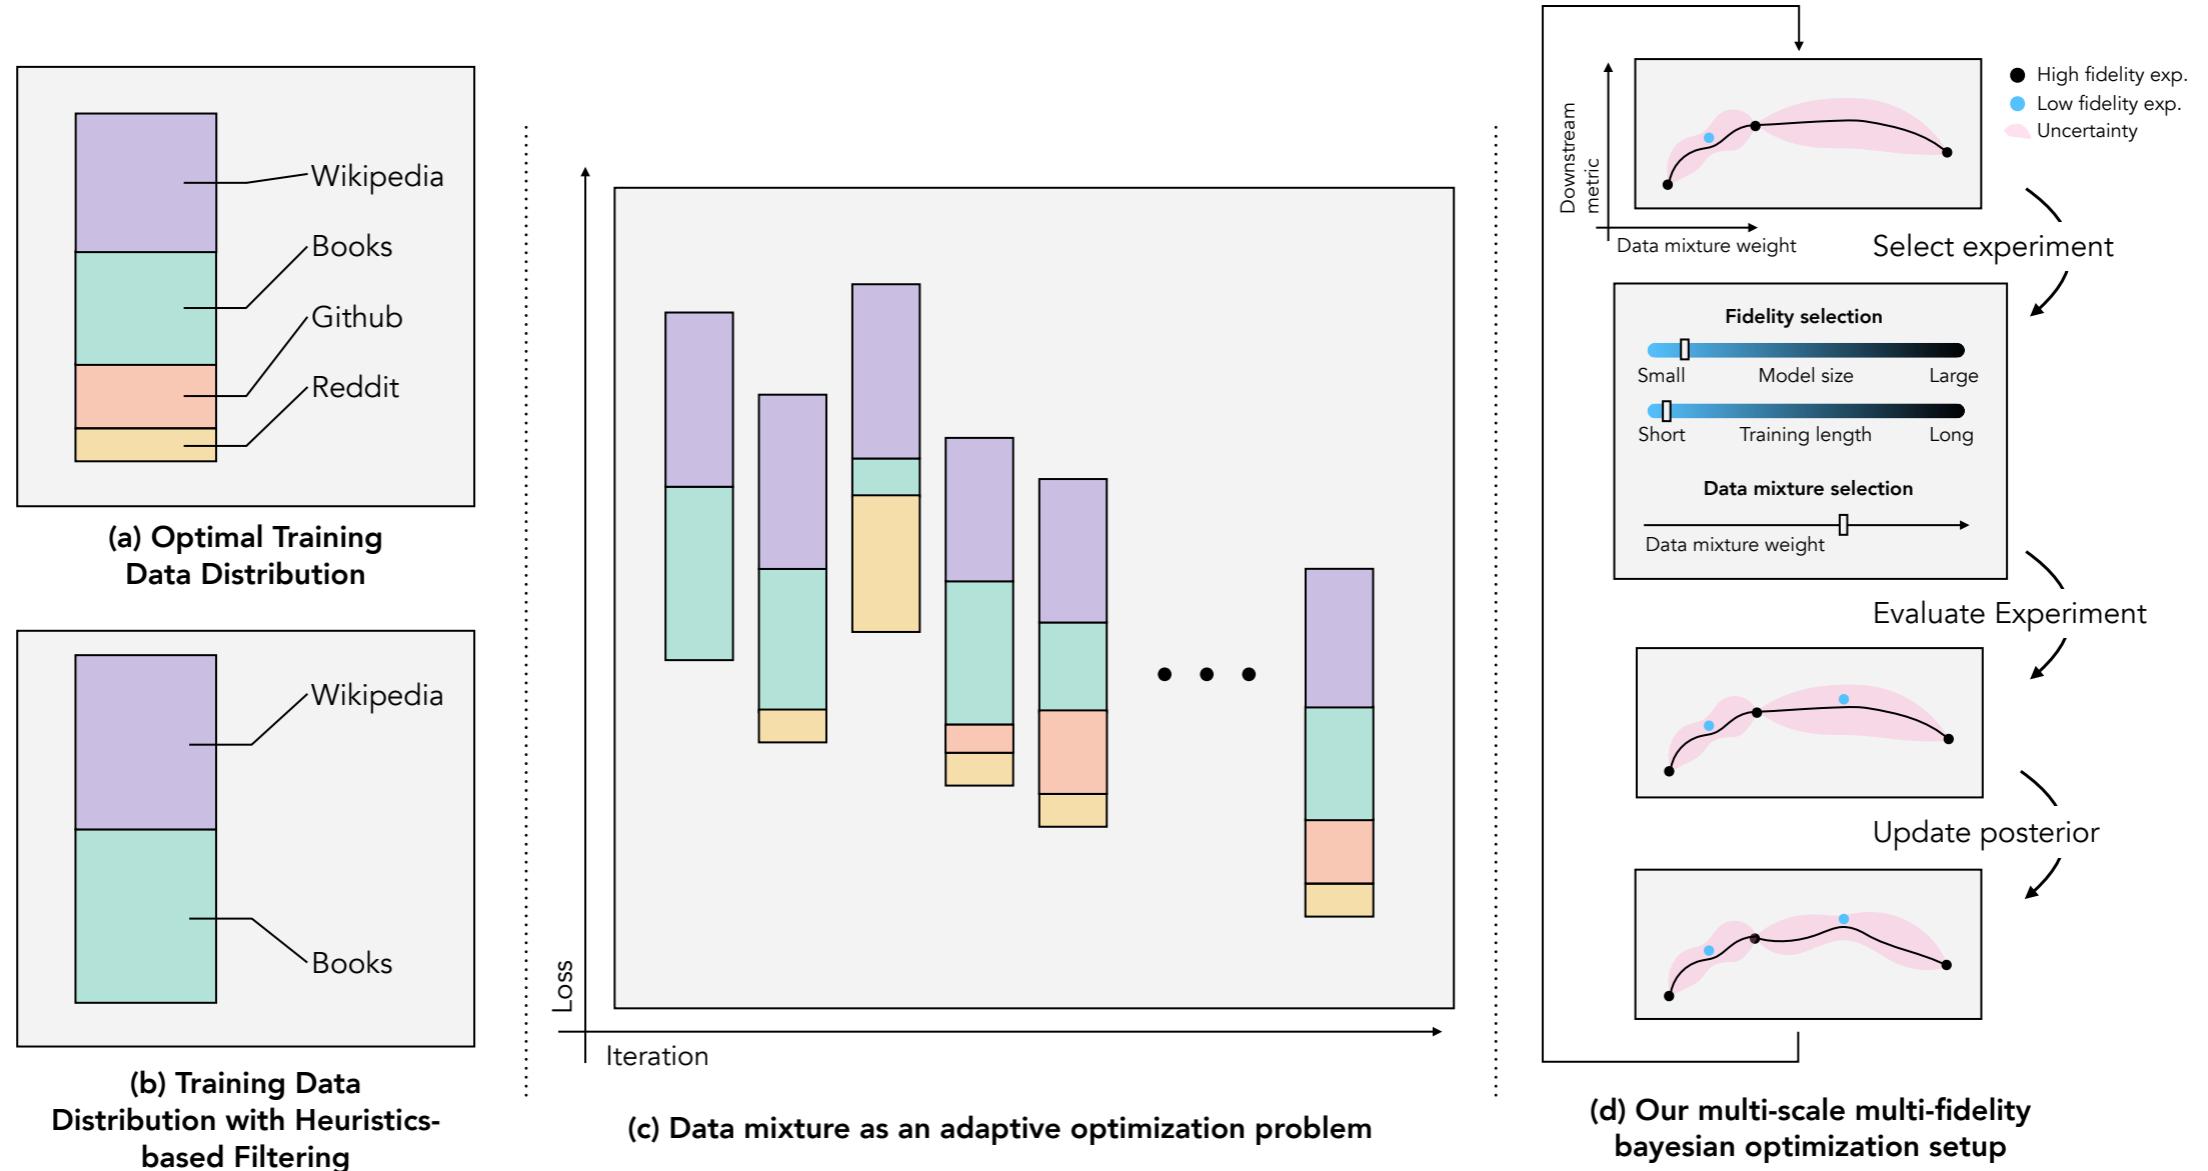
$$x^* = \arg \min f(x, s^*, m^*)$$

e.g. I want to train the best 1B model for  $s^{*}=100$  steps

I can train any  $x$  with  $s < s^*$  steps and any smaller model  $m < 1B$

The cost is  $c(s, m)$

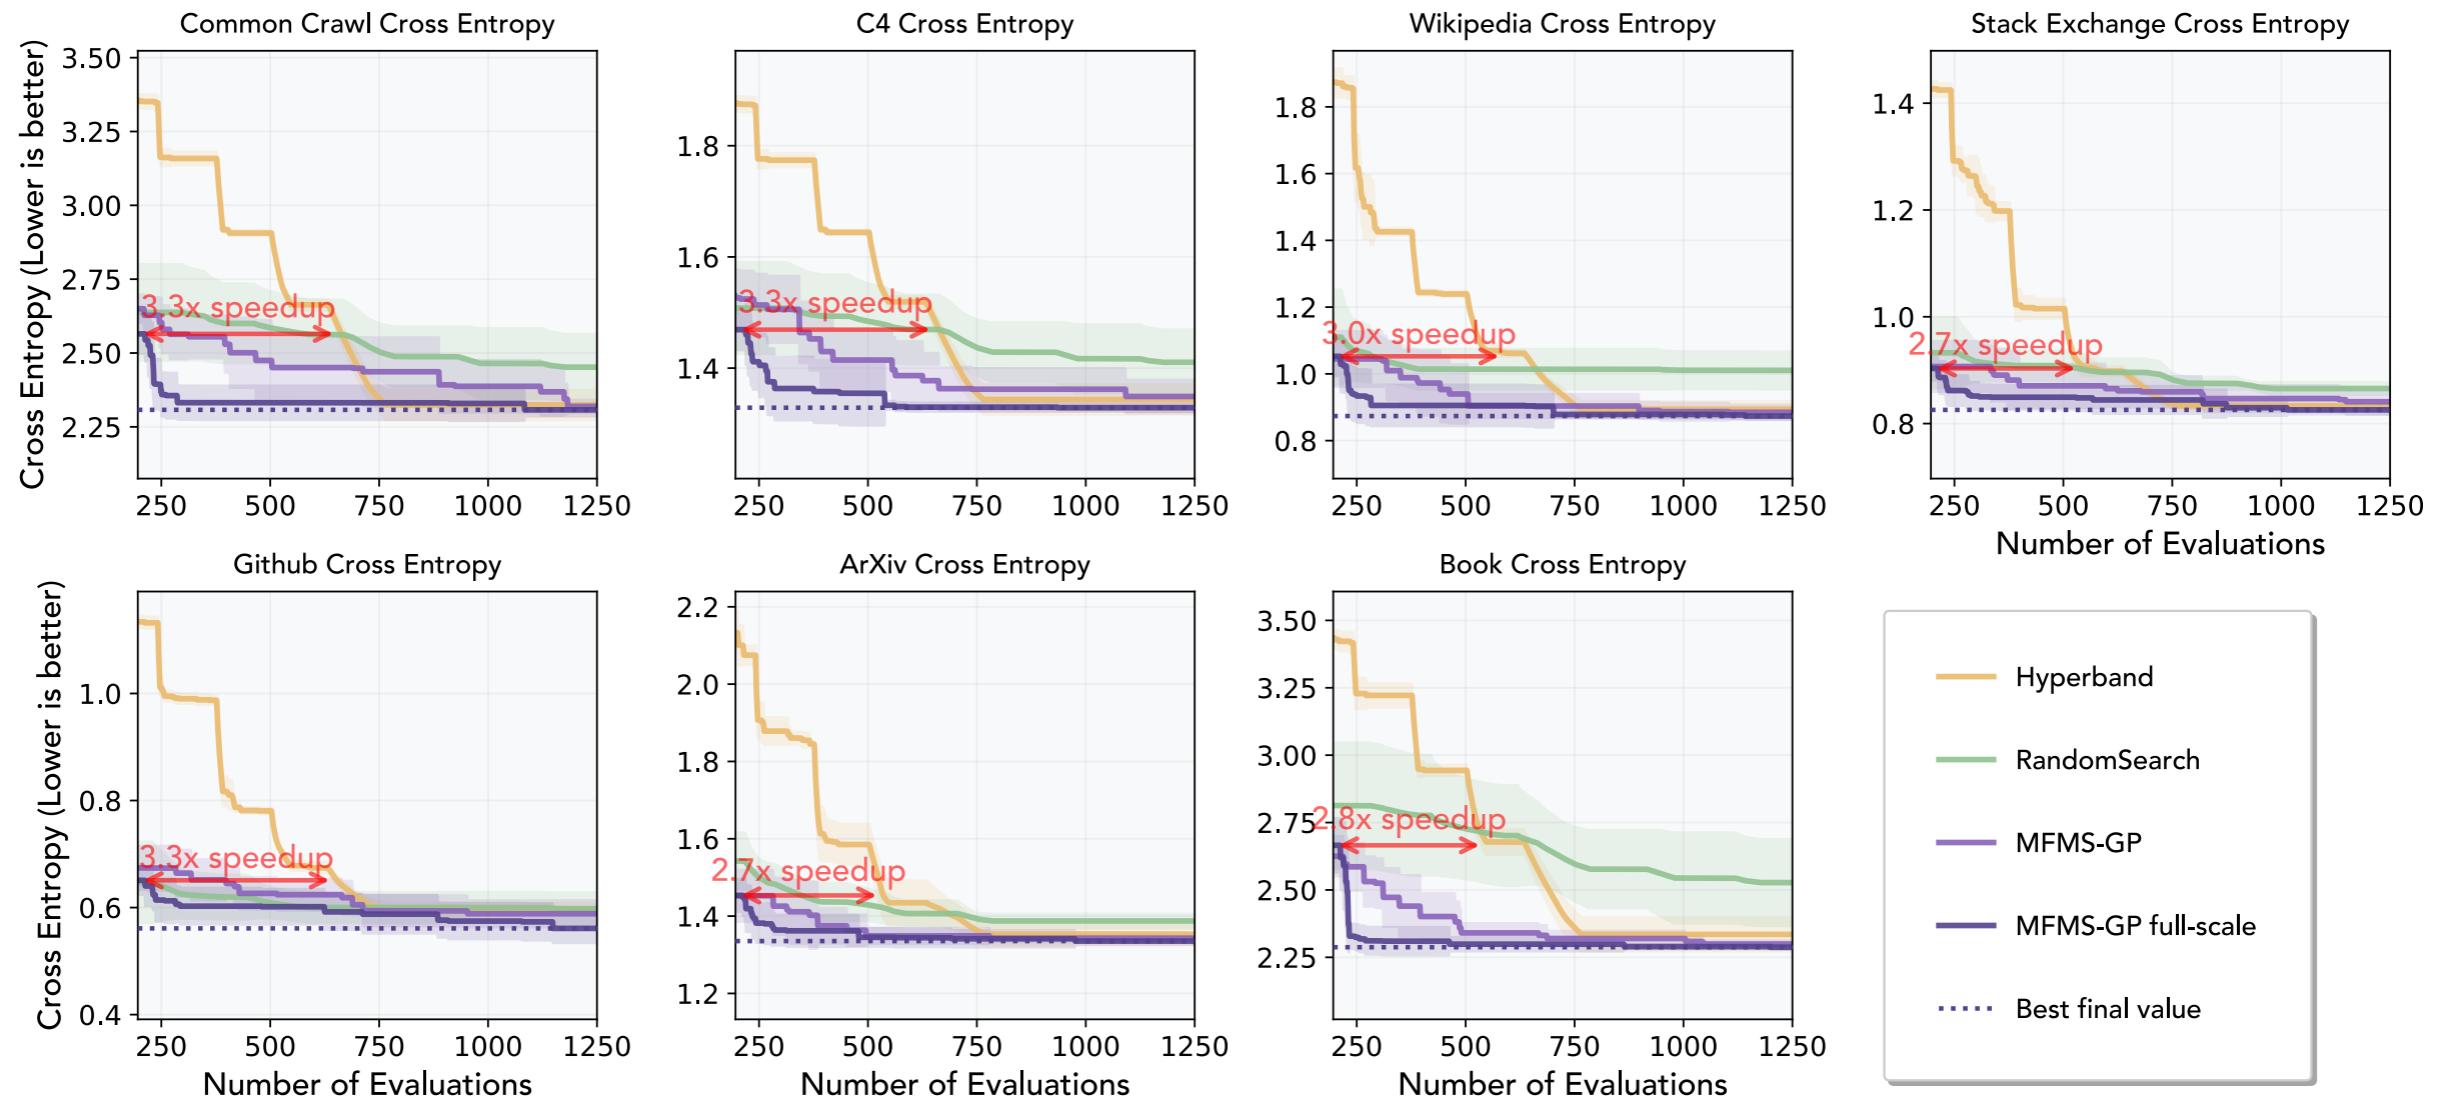
# Multi-fidelity Multi-scale Bayesian Optimization (Data Mixing)



At any iteration, I can train with data mixture  $x$ , model scale  $m$ , steps  $s$

# Multi-fidelity Multi-scale Bayesian Optimization (Data Mixing)

Extremely simple implementation of GP (using EI) works

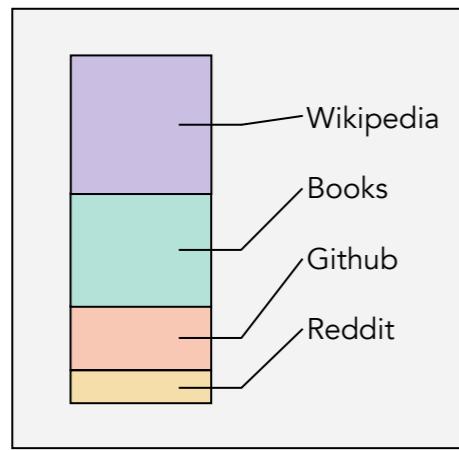


# Recap on Bayesopt

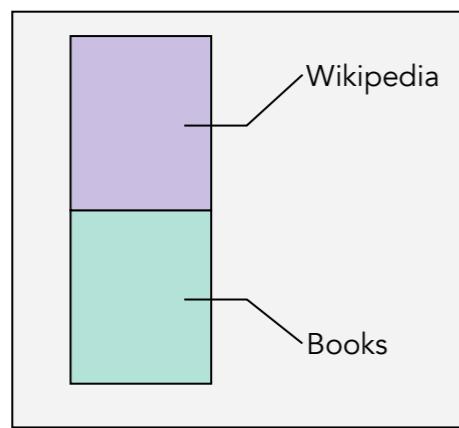
- Instead of believing in a line, we can use Bayesian optimization
- Data mixing coefficients (could) transfer better from smaller scale experiments
- Difference between model scale and steps provide rich structure for future work

# Something we've been missing: change data mixture on the fly

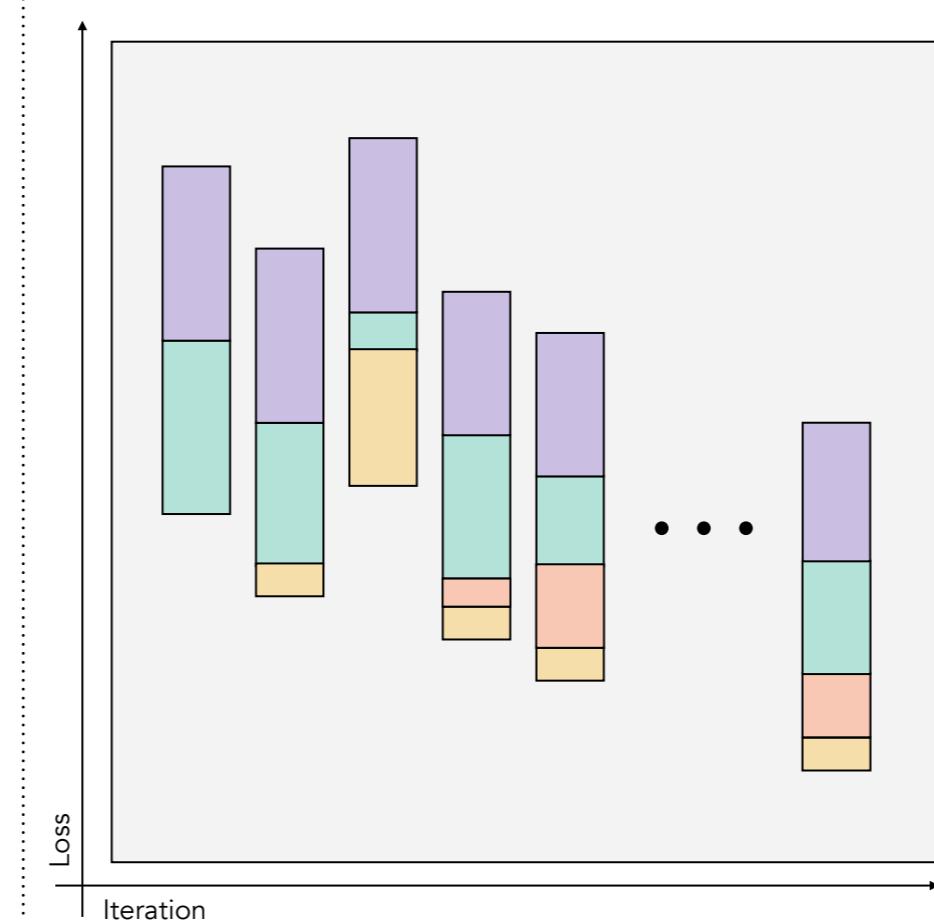
Akin to Curriculum learning: learn easy then hard



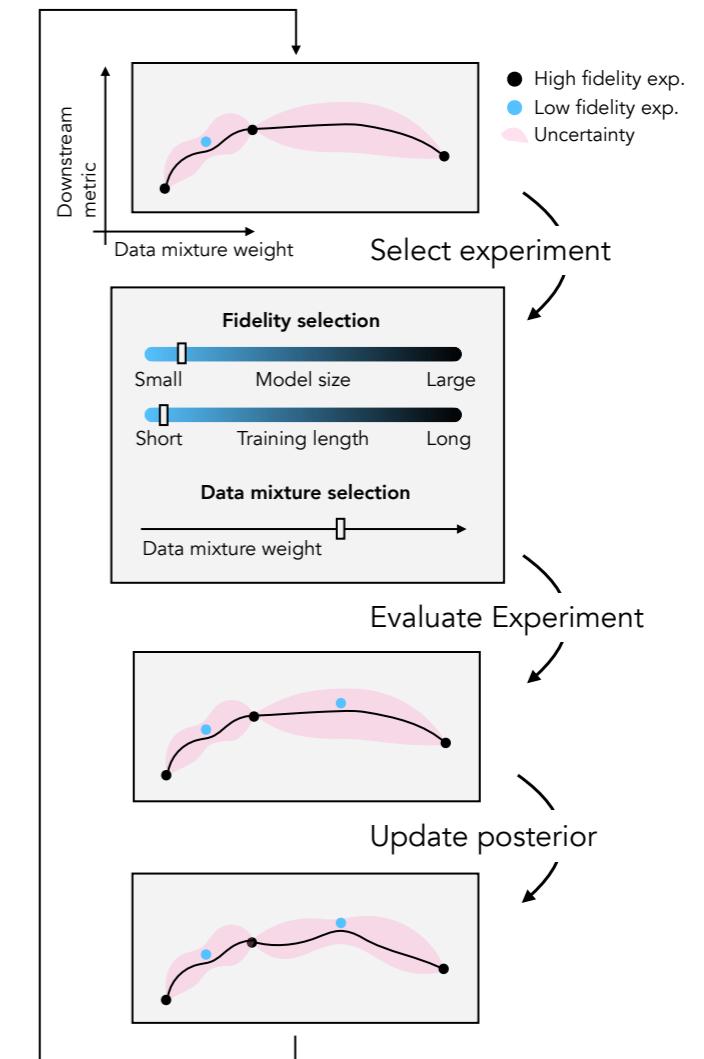
(a) Optimal Training Data Distribution



(b) Training Data Distribution with Heuristics-based Filtering



(c) Data mixture as an adaptive optimization problem

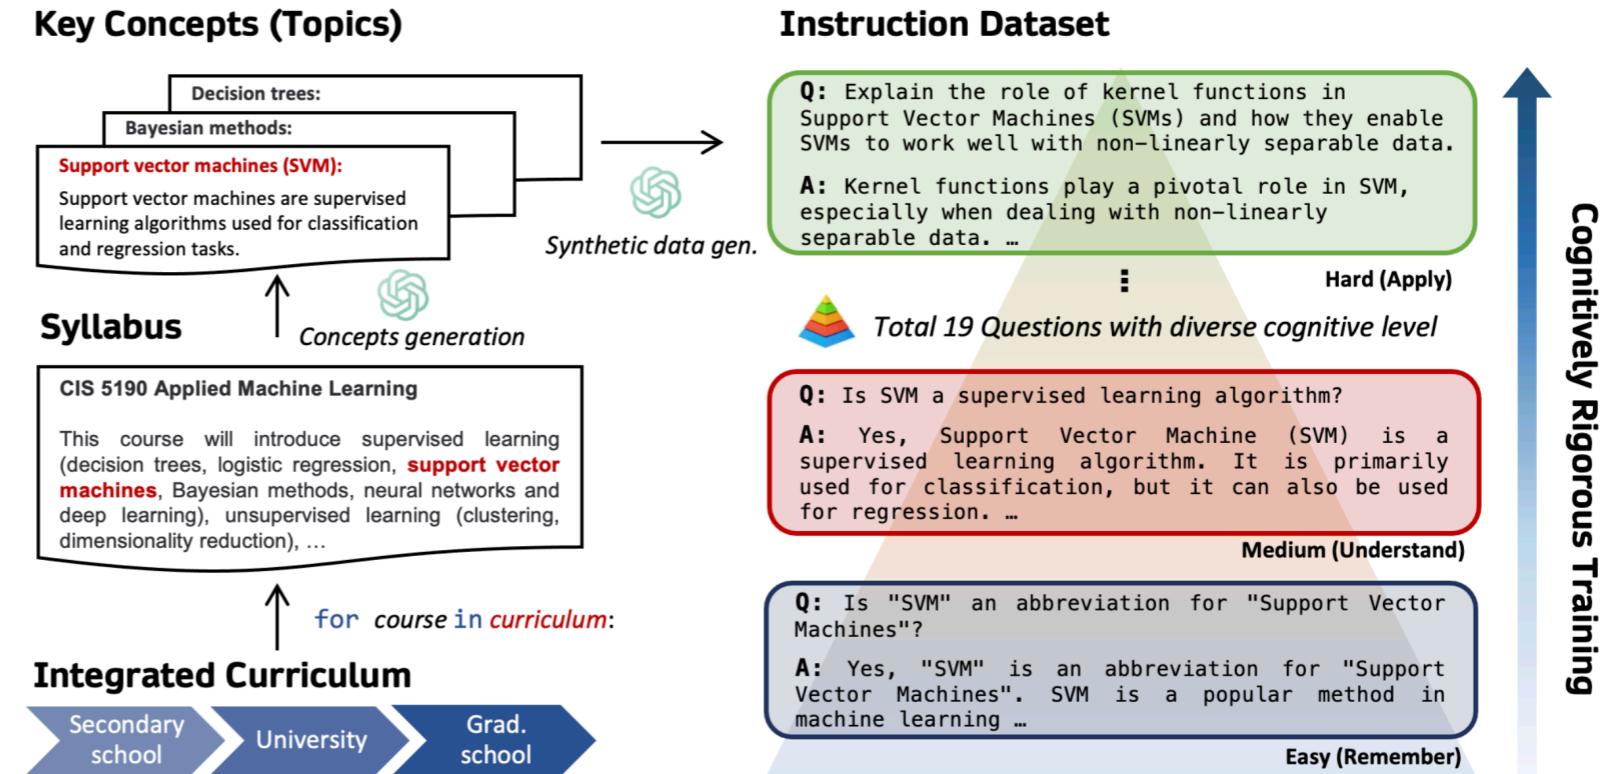


(d) Our multi-scale multi-fidelity bayesian optimization setup

# Some heuristic curriculum training

Lee et al., 2023, <https://arxiv.org/pdf/2310.09518>

## Human Curriculum



## Sequence-length

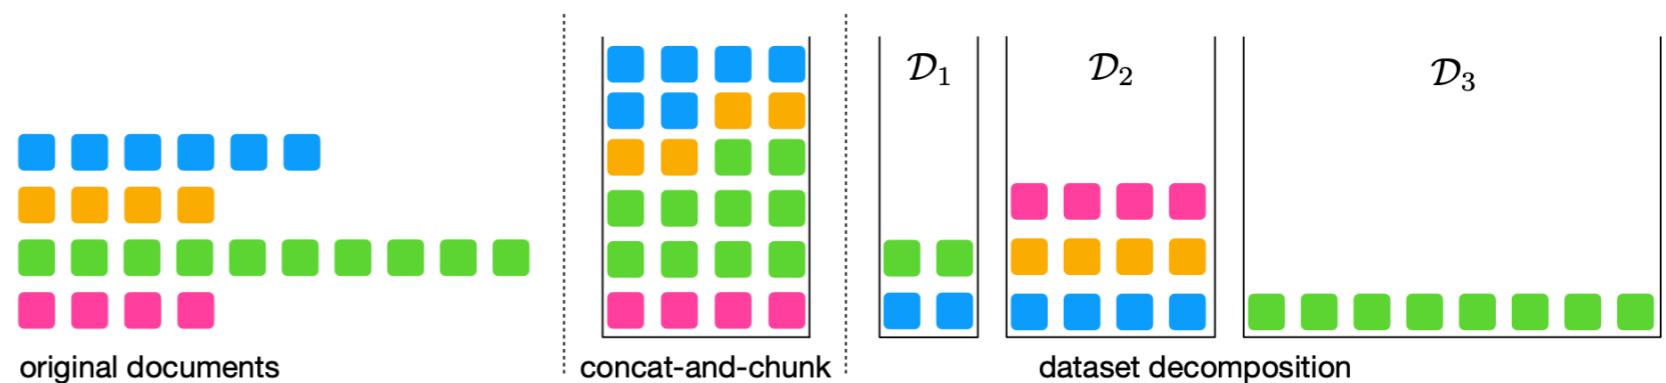


Figure 2: Each cell in the figure represents a token. **Left:** Original documents with variable lengths. **Middle:** Concat-and-chunk baseline to form sequences with a fixed target length (here = 4). **Right:** Dataset decomposition method with  $\mathcal{D}_1$ ,  $\mathcal{D}_2$ , and  $\mathcal{D}_3$  buckets .

Pouransari et al., 2024, <https://arxiv.org/pdf/2405.13226>

# Training-dynamic-based approaches

Train more on higher loss

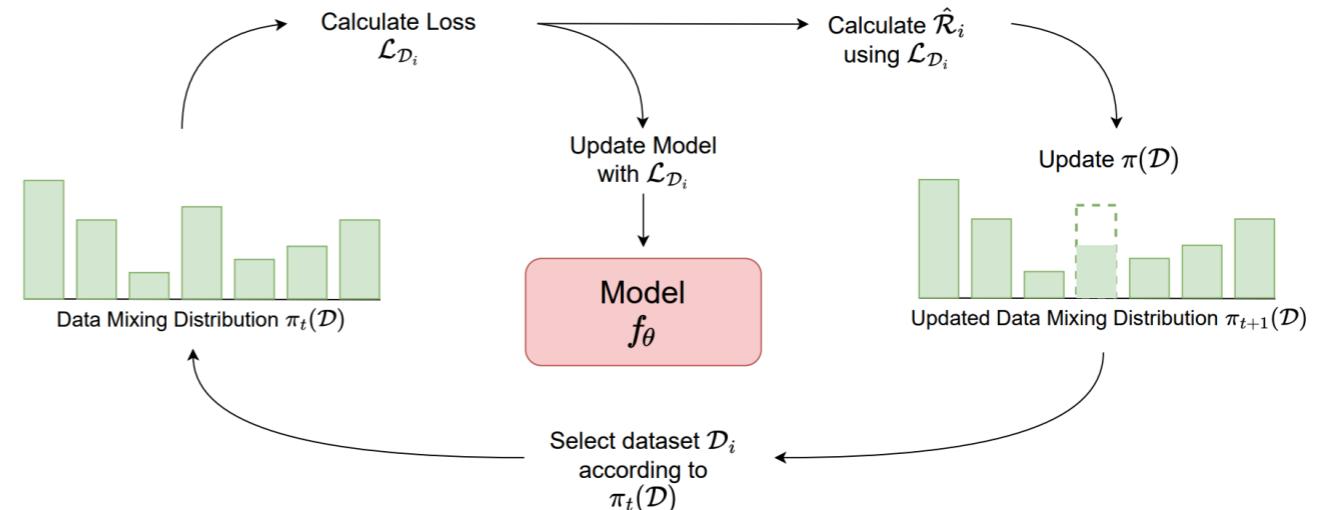


Figure 2: **Overview of Online Data Mixing (ODM) as a multi-armed bandit.** At each iteration of training,  $t$ , a dataset  $\mathcal{D}_i$  is sampled according to the data mixing distribution  $\pi$ . The loss  $\mathcal{L}_{\mathcal{D}_i}$  is calculated w.r.t the model  $f_\theta$  and subsequently used to update the model. Simultaneously, a reward  $\hat{\mathcal{R}}_i$  is calculated and used to update  $\pi$  for the next iteration,  $i + 1$ .

Albalak et al., 2024, <https://arxiv.org/pdf/2312.02406>

Use (domain-specific)  
Scaling Law to tell which  
is more learnable

$$\frac{d\widehat{\mathcal{L}}_k(n)}{dn} = \frac{-\alpha_k \beta_k n^{-\alpha_k}}{n} = -\frac{1}{n} \quad \begin{matrix} \alpha_k \\ \text{Learning speed} \end{matrix} \quad \begin{matrix} (\widehat{\mathcal{L}}_k(n) - \varepsilon_k) \\ \text{Reducible loss} \end{matrix} .$$

Jiang et al., 2024, <https://arxiv.org/pdf/2410.11820>

# Training-dynamic-based approaches

$$U^{(t)}(S; z^{(\text{val})}) := \ell(w_t, z^{(\text{val})}) - \ell(\tilde{w}_{t+1}(S), z^{(\text{val})})$$

(which data improves validation loss the most)

Expensive to compute.

Applicable only for selecting a mini batch from a batch

## **Recap on adaptive data mixing:**

- Ongoing research. Extremely simple heuristics at the moment
- Likely a lot to do here

# Data Categories

Coarse categories

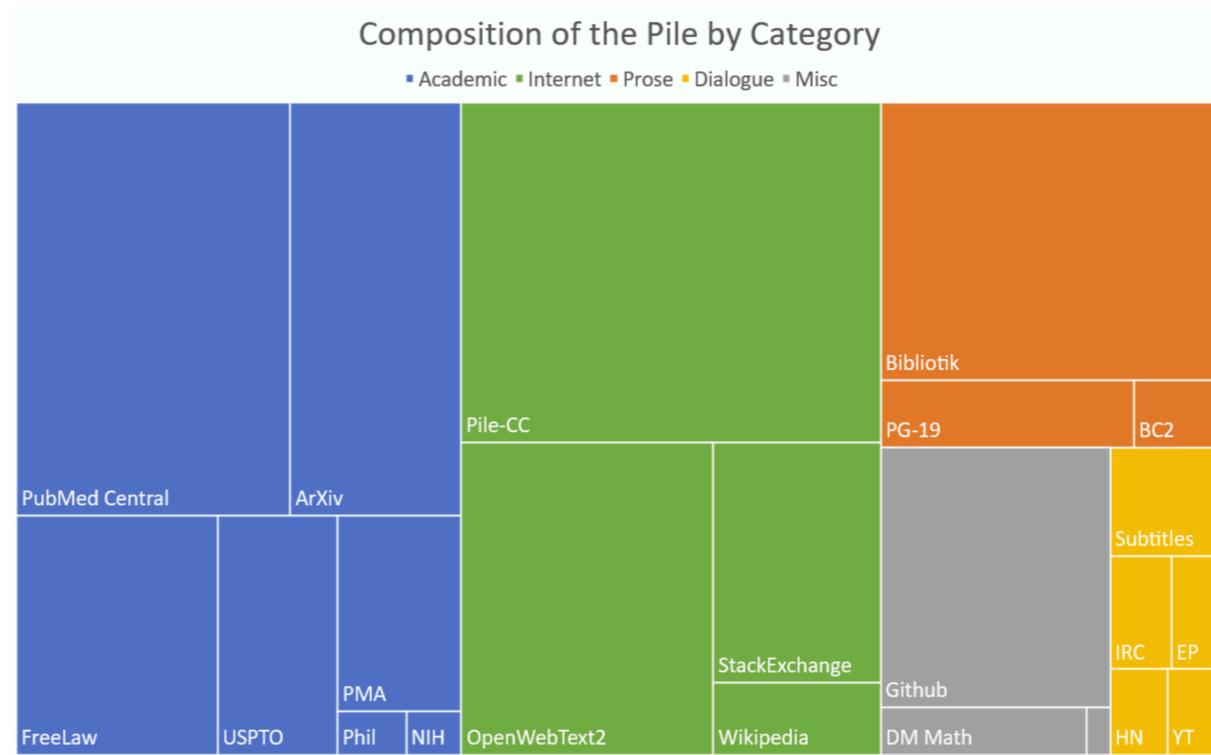


Figure 1: Treemap of Pile components by effective size.

Is there anything better?