Hongseok Namkoong

Columbia Business School
Kravis 933, 625 W 130th Street
New York, NY 10027

Email: namkoong@gsb.columbia.edu
Homepage: http://hsnamkoong.github.io
Google Scholar: https://bit.ly/3hmDmjY

Employment

Assistant Professor, Decision, Risk, and Operations Division, Columbia Business School, 2020—Present

LinkedIn Scholar, Trust & Responsible AI, LinkedIn, 2022–Present

Research Scientist, Core Data Science, Facebook, 2019–2020

Research Assistant, Peter W. Glynn and John C. Duchi, Stanford University, 2014–2019

Education

Ph.D. Management Science and Engineering, Stanford University, 2019
Advisors: John C. Duchi and Peter W. Glynn

M.S. Statistics, Stanford University, 2017

B.S. Summa Cum Laude. Industrial Engineering and Mathematics, KAIST, 2013

Teaching

B8103: Business Analytics II (MBA, MS), Columbia University, 2021-2025

B9145: Reliable Statistical Learning (PhD), Columbia University, 2020, 2023, 2025

B8784: Business and Society (MBA), Guest Instructor, Columbia University, 2024-25

Papers under review

Second+ round

P1 E. Che and H. Namkoong. Adaptive experimentation at scale: A computational framework for flexible batches. *arXiv:2303.11582 [cs.LG]*, 2023. Under second round review in Operations Research.

P2 T. Cai, H. Namkoong, and S. Yadlowsky. Diagnosing model performance under distribution shift. *arXiv:2303.02011 [stat.ML]*, 2023. Under third round review in Operations Research; Conference version appeared Symposium on Foundations of Responsible Computing 2023

P3 S. Jeong and H. Namkoong. Robust causal inference under covariate shift via worst-case subpopulation treatment effect. *arXiv:2007.02411 [stat.ML]*, 2022. Under second round review in Management Science. Short version appeared in Conference on Learning Theory 2020.

P4 H. Namkoong, S. Daulton, and E. Bakshy. Distilled Thompson sampling: Practical and efficient Thompson sampling via imitation learning. *arXiv:2011.14266 [cs.LG]*, 2024. Under second round review in Manufacturing & Service Operations Management. Selected for an oral presentation at the Neurips 2020 OfflineRL Workshop.

P5 J. Liu*, T. Wang*, P. Cui, and H. Namkoong. Rethinking distribution shifts: Empirical analysis and inductive modeling for tabular data. *arXiv:2307.05284 [cs.LG]*, 2024. Major revision in Management Science; Conference version appeared in NeurIPS 2023.

P6 E. Che, J. Dong, and H. Namkoong. Differentiable discrete event simulation for queuing network control. *arXiv:2409.03740 [cs.LG]*, 2024. Major revision in Operations Research.

P7 J. Lee, H. Namkoong, and Y. Zeng. Design and scheduling of an AI-based queueing system. *arXiv:2406.06855 [math.OC]*, 2024. Major revision in Management Science; Selected for presentation at SIG Day 2025.

First round

P8 J. Liu, T. Wang, H. Lam, H. Namkoong, and J. Blanchet. DRO: A python library for distributionally robust optimization in machine learning. *arXiv:2505.23565 [cs.LG]*, 2025.

P9 D. Mittal*, Y. Ma*, S. Joshi, and H. Namkoong. A planning framework for adaptive labeling. *arXiv:2502.06076 [cs.LG]*, 2025. Conference version appeared in NeurIPS 2024.

P10 T. Cai, H. Namkoong, D. Russo, and K. Zhang. Active exploration via autoregressive generation of missing data. *arXiv:2405.19466 [cs.LG]*, 2024. Selected for presentation at the Econometric Society Interdisciplinary Frontiers: Economics and AI+ML conference.

P11 E. Che, D. Jiang, H. Namkoong, and J. Wang. Optimization-driven adaptive experimentation. *arXiv:2408.04570 [cs.LG]*, 2024. Selected for oral presentations at the Econometric Society Interdisciplinary Frontiers: Economics and AI+ML conference and Conference on Digital Experimentation.

P12 N. Ye and H. Namkoong. Exchangeable sequence models quantify uncertainty over latent concepts. *arXiv:2408.03307 [stat.ML]*, 2024.

P13 T. Cai*, Y. Fonseca*, K. Hou, and H. Namkoong. Constrained learning for causal inference and semiparametric statistics. *arXiv:2405.09493 [stat.ML]*, 2024.

P14 M. Li, H. Namkoong, and S. Xia. Evaluating model performance under worst-case subpopulations. *arXiv:2407.01316 [cs.LG]*, 2025. Short version appeared in NeurIPS 2021.

ML conferences

P15 T. Yen*, A. W. T. Siah*, H. Chen, T. Peng, D. Guetta, and H. Namkoong. Data mixture optimization: A multi-fidelity multi-scale bayesian framework. *arXiv:2503.21023 [cs.LG]*, 2025.

P16 A. Li, H. Chen, H. Namkoong, and T. Peng. LLM generated persona is a promise with a catch. *arXiv:2503.16527 [cs.CL]*, 2025.

P17 D. Mittal*, A. Li*, T.-C. Yen*, D. Guetta, and H. Namkoong. Architectural and inferential inductive biases for exchangeable sequence modeling. *arXiv:2503.01215 [cs.LG]*, 2025.

P18 K. Zhang, T. Cai, H. Namkoong, and D. Russo. Contextual Thompson sampling via generation of missing data. *arXiv:2502.07064 [cs.LG]*, 2025.

P19 Y. Zeng*, J. Liu*, H. Lam, and H. Namkoong. LLM embeddings improve test-time adaptation to tabular $Y|X$-shifts. *arXiv:2410.07395 [cs.LG]*, 2024.

Working papers not yet submitted

P20 J. Wang, E. Che, D. Jiang, and H. Namkoong. AExGym: Benchmarks and environments for adaptive experimentation. *arXiv:2408.04531 [cs.LG]*, 2024.

P21 H. Namkoong, Y. Xu, and S. Zheng. Triply robust causal estimation for continuous treatments. *Working Paper*, 2025.

P22 D. Mittal, S. Zheng, J. Dong, and H. Namkoong. Data-driven stochastic modeling using autoregressive sequence models. *Working Paper*, 2025.

P23 Y. Ma and H. Namkoong. A practical minimax approach to causal inference with limited overlap. *Working Paper*, 2025. Conference version presented in CODE 2024.

P24 A. Boyarsky, N. Egami, and H. Namkoong. Sensitivity analysis for external validity. *Working Paper*, 2025. Conference version presented in ACIC 2024.

Publications

Journal publications

J1 H. Namkoong*, Y. Ma*, and P. W. Glynn. Minimax optimal estimation of stability under distribution shift. *Operations Research*, 2025.

J2 S. Yadlowsky, H. Namkoong, S. Basu, J. Duchi, and L. Tian. Bounds on the conditional and average treatment effect with unobserved confounding factors. *Annals of Statistics*, 50(5):2587–2615, 2022.

J3 J. C. Duchi, T. Hashimoto, and H. Namkoong. Distributionally robust losses against mixture covariate shifts. *Operations Research*, 2022.

J4 J. C. Duchi and H. Namkoong. Learning models with uniform performance via distributionally robust optimization. *Annals of Statistics*, 49(3):1378–1406, 2021.

J5 J. C. Duchi, P. W. Glynn, and H. Namkoong. Statistics of robust optimization: A generalized empirical likelihood approach. *Mathematics of Operations Research*, 46(3):946–969, 2021.

J6 J. C. Duchi and H. Namkoong. Variance-based regularization with convex objectives. *Journal of Machine Learning Research*, 2019.

Refereed conference proceedings[1]

C1 J. Wang, T. Zollo, R. Zemel, and H. Namkoong. Adaptive elicitation of latent information using natural language. In *Proceedings of the 42nd International Conference on Machine Learning*, 2025. Selected for spotlight presentation at ICLR 2025 Workshop on Quantify Uncertainty and Hallucination in Foundation Models.

C2 T. Zollo*, A. Siah*, N. Ye, A. Li, and H. Namkoong. PersonalLLM: Tailoring LLMs to individual preferences. In *Proceedings of the Thirteenth International Conference on Learning Representations*, 2024.

C3 B. Hsu, C. DiCiccio, N. Pillai, and H. Namkoong. From models to systems: A comprehensive framework for ai system fairness in compositional recommender systems. In *Proceedings of the Algorithmic Fairness Through the Lens of Metrics and Evaluation*, volume 279 of *Proceedings of Machine Learning Research*, pages 8–37. PMLR, 2025.

C4 D. Mittal*, Y. Ma*, S. Joshi, and H. Namkoong. Adaptive labeling for efficient out-of-distribution model evaluation. In *Advances in Neural Information Processing Systems 37*, 2024.

C5 H. Chen, A. Li, E. Che, T. Peng, J. Dong, and H. Namkoong. QGym: Scalable simulation and benchmarking of queuing network controllers. In *Advances in Neural Information Processing Systems 37, Datasets and Benchmark Track*, 2024.

---

[1]Papers displayed in gray are superseded by long versions.

C6 J. Liu*, T. Wang*, P. Cui, and H. Namkoong. On the need for a language describing distribution shifts: Illustrations on tabular datasets. In *Advances in Neural Information Processing Systems 36*, 2023.

C7 A. Boyarsky, H. Namkoong, and J. Pouget-Abadie. Modeling interference via experiment rollout. In *Proceedings of the 24th ACM conference on Economics and Computation*, 2023.

C8 T. Cai, H. Namkoong, and S. Yadlowsky. Diagnosing model performance under distribution shift. In *Symposium on Foundations of Responsible Computing*, 2023.

C9 M. Wortsman, G. Ilharco, S. Y. Gadre, R. Roelofs, R. Gontijo-Lopes, A. S. Morcos, H. Namkoong, A. Farhadi, Y. Carmon, S. Kornblith, and L. Schmidt. Model soups: averaging weights of multiple fine-tuned models improves accuracy without increasing inference time. In *Proceedings of the 39th International Conference on Machine Learning*, 2022

C10 M. Wortsman, G. Ilharco, J. W. Kim, M. Li, S. Kornblith, R. Roelofs, R. Gontijo-Lopes, H. Hajishirzi, A. Farhadi, H. Namkoong, and L. Schmidt. Robust fine-tuning of zero-shot models. In *Proceedings of the 32nd IEEE Conference on Computer Vision and Pattern Recognition*, 2022. Selected for a full oral presentation

C11 M. Li, H. Namkoong, and S. Xia. Evaluating model performance under worst-case subpopulations. In *Advances in Neural Information Processing Systems 34*, 2021.

C12 H. Namkoong*, R. Keramati*, S. Yadlowsky*, and E. Brunskill. Off-policy policy evaluation for sequential decisions under unobserved confounding. In *Advances in Neural Information Processing Systems 33*, 2020.

C13 S. Jeong and H. Namkoong. Robust causal inference under covariate shift via worst-case subpopulation treatment effect. In *Conference on Learning Theory*, 2020.

C14 M. O'Kelly*, A. Sinha*, H. Namkoong*, J. Duchi, and R. Tedrake. Scalable end-to-end autonomous vehicle testing via rare-event simulation. In *Advances in Neural Information Processing Systems 31*, 2018.

C15 R. Volpi*, H. Namkoong*, J. Duchi, V. Murino, and S. Savarese. Generalizing to unseen domains via adversarial data augmentation. In *Advances in Neural Information Processing Systems 31*, 2018.

C16 T. Hashimoto, M. Srivastava, H. Namkoong, and P. Liang. Fairness without demographics in repeated loss minimization. In *International Conference on Machine Learning*, 2018. Best Paper Runner-up Award.

C17 A. Sinha*, H. Namkoong*, and J. Duchi. Certifiable distributional robustness with principled adversarial training. In *International Conference on Learning Representations*, 2018. Selected for a full oral presentation; 2% of submissions.

C18 H. Namkoong and J. C. Duchi. Variance regularization with convex objectives. In *Advances in Neural Information Processing Systems 30*, 2017. Best Paper Award.

C19 H. Namkoong, A. Sinha, S. Yadlowsky, and J. C. Duchi. Adaptive sampling probabilities for non-smooth optimization. In *International Conference on Machine Learning*, pages 2574–2583, 2017.

C20 H. Namkoong and J. C. Duchi. Stochastic gradient methods for distributionally robust optimization with f-divergences. In *Advances in Neural Information Processing Systems 29*, 2016.

Technical reports & software

T1 G. Ilharco, M. Wortsman, N. Carlini, R. Taori, A. Dave, V. Shankar, H. Namkoong, J. Miller, H. Hajishirzi, A. Farhadi, and L. Schmidt. Openclip: Open-sourced implementation of clip, Jul 2021. URL https://doi.org/10.5281/zenodo.5143773

T2 A. Sinha*, H. Namkoong*, R. Volpi, and J. Duchi. Certifying some distributional robustness with principled adversarial training. *Technical Report*, 2020.

T3 H. Namkoong, J. C. Duchi, and P. W. Glynn. Proofs for empirical likelihood with general f-divergences. *Technical Report*, 2018.

Honors & Awards

Interdisciplinary Seed Grant Program Competition Award, Columbia University, 2024

Amazon Research Award, 2022

Best Paper Finalist for "Robust fine-tuning of zero-shot models", *Conference on Computer Vision and Pattern Recognition (CVPR)*, 2022

Best Student Paper Award for "Statistics of Robust Optimization: A Generalized Empirical Likelihood Approach", *INFORMS Applied Probability Society*, 2018

Best Paper Runner Up Award for "Fairness Without Demographics in Repeated Loss Minimization", *International Conference on Machine Learning (ICML)*, 2018

Best Paper Award for "Variance Based Regularization with Convex Objectives", *Neural Information Processing Systems (NeurIPS)*, 2017

Samsung Fellowship, 2013–2018

Department Fellowship, Management Science and Engineering, Stanford, 2013–2018

KAIST President's Award (graduated top of class in the School of Engineering), 2013

Undergraduate Research Award, First Place, Department of Industrial and Systems Engineering, 2012

Professional Service

Senior Program Committee

2021-25	Area chair, NeurIPS
2023-25	Area chair, ICML

Reviewing

Journals *Operations Research, Management Science, Journal of the American Statistical Association, Journal of the Royal Statistical Society: Series B, Mathematical Programming, SIAM Journal on Mathematics of Data Science, Journal of Machine Learning Research, Transactions on Pattern Analysis and Machine Intelligence, Automatica.*

Conferences *Neural Information Processing Systems, International Conference on Machine Learning, Conference on Learning Theory, Conference on Algorithmic Learning Theory, MSOM Special Interest Group*

Workshop Organization

2025	Co-organizer, Columbia Catalyst Workshop, Data Science Institute
2024	Co-Organizer, four INFORMS invited sessions on AI-driven decision-making
2024	Co-organizer, ACIC Invited Session on "External Validity"
2023	Co-organizer, NeurIPS workshop on "Distribution shifts (DistShift)"
2023	Co-Organizer, three INFORMS invited sessions on experimentation
2023	Organizer, Digital Future Initiative workshop on "Challenges in Operationalizing Responsible AI"
2023	Organizer, INFORMS APS invited session "Frontiers in Sequential Learning"
2022	Co-organizer, NeurIPS workshop on "Distribution shifts: connecting methods and applications (DistShift)"
2022	Organizer, INFORMS invited session "Causality & Robustness"
2022	Co-organizer, INFORMS invited session "Experimentation Design"
2022	INFORMS HAS Pierskalla Award committee member
2022	Co-organizer, ICDM workshop on "Algorithms Towards Ethical and Privacy Challenges in Social Media Recommendation Systems (AESM)"
2021	Organizer, mentoring program for PhD students at the Neurips workshop on "Distribution shifts: connecting methods and applications (DistShift)"
2021	Co-organizer, NeurIPS workshop on "Distribution shifts: connecting methods and applications (DistShift)"
2021	Co-organizer, JSM invited session on "distributional robustness, validity, causality, and generalizability"
2019	Co-organizer, INFORMS invited session on "AI and machine learning"

Academic Committees

Ph.D. Coordinator, Columbia University, 2025–2026

Member, Ph.D. Admissions Committee, Columbia University, 2020–current

Member, Faculty Hiring Committee, Columbia University, 2021-22, 2023-24

Member, Ph.D. Program Review Committee, Columbia University, 2021

Public Speaking For Columbia

Panelist, CBS Alumni Club of New York, 2025

Guest Instructor with Lan Guan (CAIO at Accenture), Business and Society, 2025

"AI: Opportunities and Risks", CBS Alumni Event With Dean Maglaras, 2025

"AI: Opportunities and Risks", Chazen Institute for Global Business Board Meeting, 2024

Guest Instructor with Scott Silver (VP at YouTube), Business and Society, 2024

Panelist, Showcase on Large Language Models, Columbia Center of AI Technology, 2023

Moderator, Responsible AI and Content Moderation Bernstein Center for Leadership and Ethics 2023

External Service

NSF panel, 2024

Ad hoc reviewer, Methodology, Measurement, and Statistics Program, NSF, 2024

Research Supervision

Ph.D. Advising

Ari Boyarsky (DRO), Tiffany Cai (Statistics), Yuanzhe Ma (IEOR, co-advised with Garud Iyengar and Jay Sethuraman), Daksh Mittal (DRO, co-advised with Jing Dong), Naimeng Ye (DRO), Thomson Yen (DRO), Isaac Scheinfeld (DRO, co-advised with Hannah Li)

Alumni Ethan Che (Ph.D., 2025, co-advised with Jing Dong, at Meta Adaptive Experimentation), Yibo Zeng (Ph.D., 2024, co-advised with Henry Lam, at Meta Virtual Reality Labs)

Student collaborators Yuri Fonseca, Nicholas Galbraith, Yian Huang, Mike Li, Jiashuo Liu, Tianyu Wang, Shawn Xia, Yunbei Xu, Tom Zollo.

Research Scientists

Minghui Chen

Postdocs

Yanlin Qu (co-advised with Assaf Zeevi).

Alumni Kelly Zhang (postdoc, 2023-24, co-advised with Dan Russo, Assistant Professor at Imperial College London)

Undergraduates and MS students

Gilbert Yang (Columbia CS), Yaqin Chen (Sun Yat Sen CS), Melody Ma (Columbia CS),

Alumni Tony Chen (Ph.D. candidate at Princeton CS), Elise Han (Ph.D. candidate at Stanford CS), Leon Li (Ph.D. candidate at NYU CS), Jimmy Wang (Ph.D. candidate at Berkeley CS), Shunri Zheng (Ph.D. candidate at UIUC IE)

Ph.D. Thesis Committee

Samuel Deng (CS), Nicholas Galbraith (Statistics), Yuanling (Judy) Gan (DRO), Yian Huang (Statistics), Luofeng Liao (IEOR), Yookoon Park (CS), Sudeep Raja Putta (IEOR), Claudia Shi (CS), Madhu Sridharan (IEOR), Shangzhou (Shawn) Xia (DRO), Yunbei Xu (DRO), Yibo Zeng (IEOR), Junzhe Zhang (CS)

Invited Talks

2025	INFORMS Applied Probability Society Conference, Atlanta, GA
2025	SIG Day (Service Management), London Business School
2025	Workshop in Learning, London Business School
2025	Adaptive Experimentation Workshop, Meta, NYC
2025	Department of Industrial Engineering and Operations Research, University of California, Berkeley
2025	Operations and Logistics Division, University of British Columbia
2024	Domain Adaptation and Related Areas, Simons Institute
2024	Political Science Methods Workshop, UC Berkeley
2024	Central Applied Sciences, Meta

2024	TOPS Department Seminar, Stern School of Business, New York Unviersity
2024	INFORMS Annual Meeting, Seattle, WA
2024	Econometric Society Interdisciplinary Frontiers: Economics and AI+ML, Cornell
2024	Deployable RL Workshop Keynote, RL Conference, Amherst, MA
2024	International Symposium on Mathematical Programming, Montreal
2024	INFORMS Optimization Society Conference, Houston, TX
2024	Workshop on Optimal Transport: Theory to Applications, Humboldt University, Germany
2023	IMS International Conference on Statistics and Data Science, Lisbon, Portugal
2023	NeurIPS Tutorial "Modeling and Exploiting Data Heterogeneity under Distribution Shifts", New Orleans, LA
2023	Stochastic Networks and Applied Probability, and Performance Seminar
2023	Management Sciences and Information Systems, Rutgers Business School
2023	Invited Session, INFORMS Annual Meeting , Phoenix, AZ
2023	Department of Industrial Engineering and Management Science, Northwestern University
2023	Big Data and Machine Learning in Econometrics, Finance, and Statistics, University of Chicago
2023	IBM Watson, NY
2023	Joint Statistical Meetings, Toronto
2023	New England Statistics Symposium
2023	INFORMS Revenue Management & Pricing Conference, London, United Kingdom
2023	INFORMS Applied Probability Society Conference, Nancy, France
2023	Department of Management Science and Engineering, Stanford University
2023	AI Matrix Seminar, UT San Antonio
2023	Operations Research Center, MIT
2022	Conference on Digital Experimentation, Cambridge
2022	INFORMS Annual Meeting, Indianapolis
2022	IFDS Workshop on Distributional Robustness in Data Science, University of Washington
2022	International Conference on Continuous Optimization (Bethlehem, PA)
2022	Department of Industrial & Systems Engineering and Computer Science, KAIST
2022	College of Business and AI, KAIST
2022	International Conference on Econometrics and Statistics, Kyoto
2022	Data Science Lab, MIT
2022	Department of Economics, Columbia University
2022	Workshop on Foundations of Stable, Generalizable and Transferable Statistical Learning, Mathematical Sciences Research Institute
2021	Department of Statistics, Columbia University
2021	LinkedIn
2021	Workshop on "distributional robustness, validity, causality, and generalizability", Joint Statistical Meetings
2021	Empirical Inference Department, Max Planck Institute for Intelligent Systems
2021	Department of Mathematics, KAIST
2021	School of Data Science, Seoul National University
2021	Data Science Institute, Columbia University
2021	Decision Science Group, McCombs School of Business, UT Austin
2020	Samsung Advanced Institute of Technology, Seoul
2020	Google Brain, Cambridge
2020	Conference on Information Sciences and Systems
2020	American Causal Inference Conference
2020	SIAM Conference on Mathematics of Data Science
2019	Uber Marketplace and Uber Eats, San Francisco
2019	OIT Division, Graduate School of Business, Stanford University
2019	INFORMS Annual Meeting, Seattle
2019	Stitchfix, San Francisco

2019	Department of Computer Science, University of Wisconsin-Madison
2019	Department of Industrial and Systems Engineering, University of Wisconsin-Madison
2019	School of Operations Research and Industrial Engineering, Cornell Tech
2019	Machine Learning and Statistics Group, Microsoft Research New England
2019	Operations and Statistics Group, MIT Sloan School of Management
2019	Department of Operations Research and Industrial Engineering, UT Austin
2019	Machine Learning Department, Carnegie Mellon University
2019	Heinz College, Carnegie Mellon University
2019	Department of Industrial Engineering and Operations Research, Columbia University
2019	Decisions, Risk and Operations Division, Columbia Business School
2019	Department of Electrical and Computer Engineering, Purdue University
2019	Operations Management Division, Booth School of Business, University of Chicago
2019	Data Sciences and Operations, Marshall School of Business, University of Southern California
2018	Department of Industrial and Operations Engineering, University of Michigan
2018	Three invited talks, INFORMS Annual Meeting (Phoenix, AZ)
2018	Oral Presentation, International Conference on Learning Representations (Vancouver, Canada)
2017	Oral Presentation, Neural Information Processing Systems (Long Beach, CA)
2016	Department of Industrial and Systems Engineering, KAIST
2016	Young Researchers Workshop, School of ORIE, Cornell University

Outside Activities

Columbia Business School requires faculty members to disclose any activities that might present a real or apparent conflict of interest. I consult at LinkedIn, serving as a LinkedIn Scholar.